

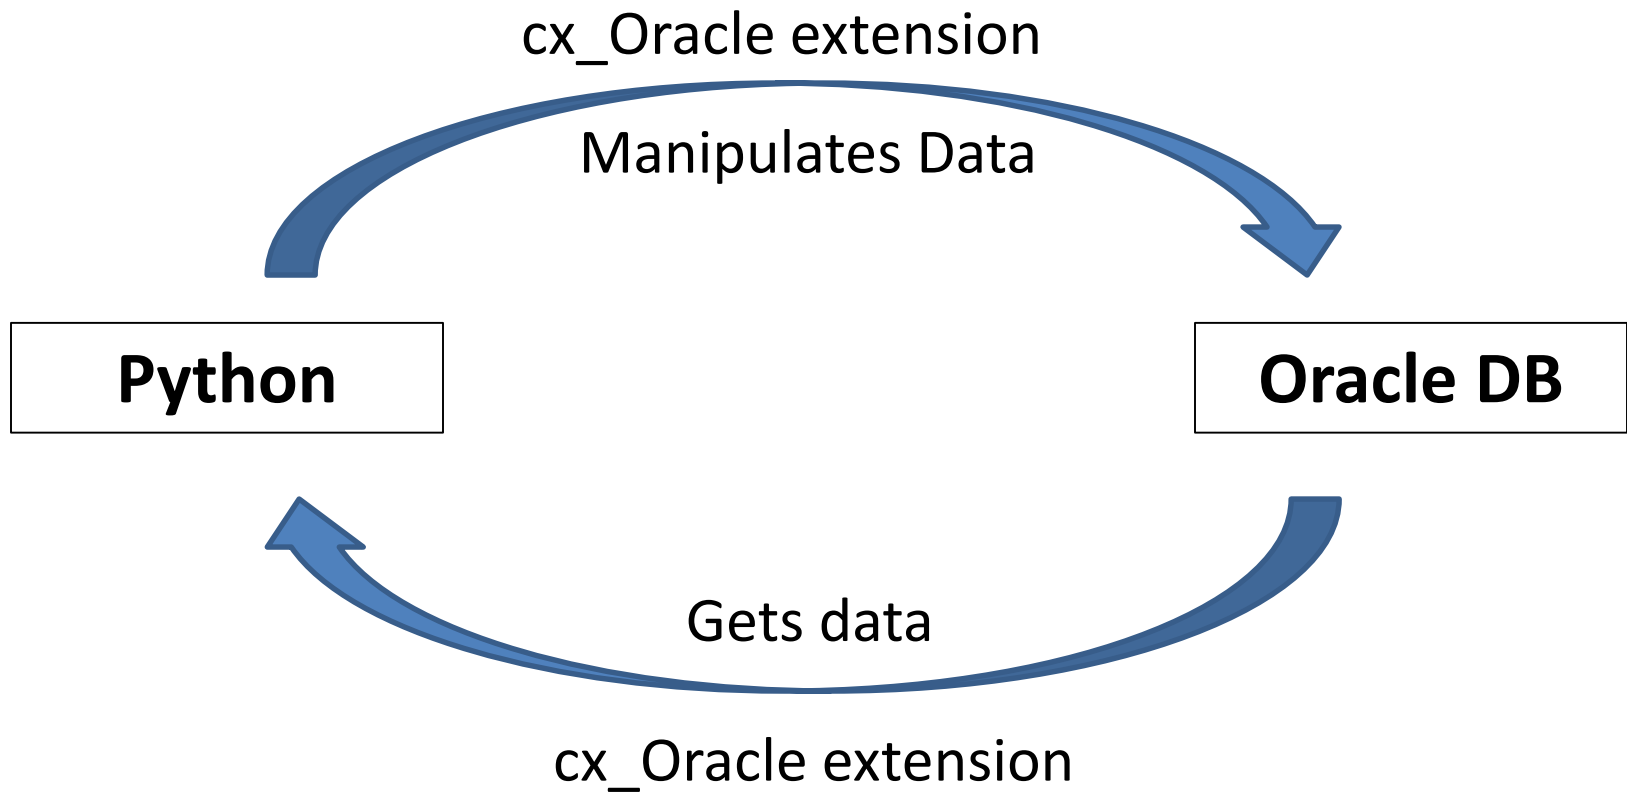
STSCI 4060

Lecture File 10

Xiaolong Yang
(xy44@cornell.edu)

Integrate Python with Oracle DB

Python and Oracle



Install cx_Oracle 7.x Extension

- **cx_Oracle** is a Python extension module for accessing Oracle databases and conforming to the Python database API specifications.
- This module is currently built against Oracle 11.2, 12.1, 12.2 and 18.3 and Python 2.7 as well as Python 3.5, 3.6 and 3.7.
- Go to https://oracle.github.io/python-cx_Oracle/ install the extension.
- For most cases, you just need to run the following command at your DOS shell:

python -m pip install cx_Oracle --upgrade

Note: in case of Spyder, if cx_Oracle is not installed, go to **Environments** and select **Not installed** and then **cx_Oracle** to install and upgrade. You may see an older version of cx_Oracle but it will work well.

The Oracle Database

- The most popular relational database system in the market/industry.
- It was introduced in STSCI 5060.
- It can integrate well with the SAS software for statistical data analysis.
- It can also integrate well with Python for general programming purposes, including data analyses.
- ...

Install Oracle 11g XE

- If you do not have a local version of Oracle 11g XE, download it from <http://www.oracle.com/technetwork/database/databases/e-technologies/express-edition/downloads/index.html> based on your computer's OS and install it.
- Also, download **Oracle SQL Developer** from the same website and install it.

Create an Oracle User Account

- While you are logged in as an Oracle **system user**, you have the privilege to create a user.
- Use **your first name** as the new database user name (if it was not been created before) and assign **a password** of your choice (write it down).
- Enter the following statement:
CREATE USER *first_name* IDENTIFIED BY *your_password*;
Then it displays “User created.”
- Grant privileges to the new user:
GRANT ALL PRIVILEGES TO *first_name*;
Then it should display “Grant succeeded.”
- Exit SQL by typing “**exit**” or “**quit**.”
- Enter “**sqlplus**” at the prompt and log on with your new user name and password.
- You can check to see who the user is by typing the “**show user**” command.

Create a Table in Your New Database

One way to create a table is based on an existing table in other databases. The following statement creates a table called DEPARTMENTS from the DEPARTMENTS table in the HR database.

```
CREATE TABLE departments AS  
SELECT * FROM hr.departments;
```

You can show your table contents with the following query:

```
SELECT * FROM departments;
```

```
SQL> select * from departments;
```

DEPARTMENT_ID	DEPARTMENT_NAME	MANAGER_ID	LOCATION_ID
10	Administration	200	1700
20	Marketing	201	1800
30	Purchasing	114	1700
40	Human Resources	203	2400
50	Shipping	121	1500
60	IT	103	1400
70	Public Relations	204	2700
80	Sales	145	2500
90	Executive	100	1700
100	Finance	108	1700
110	Accounting	205	1700

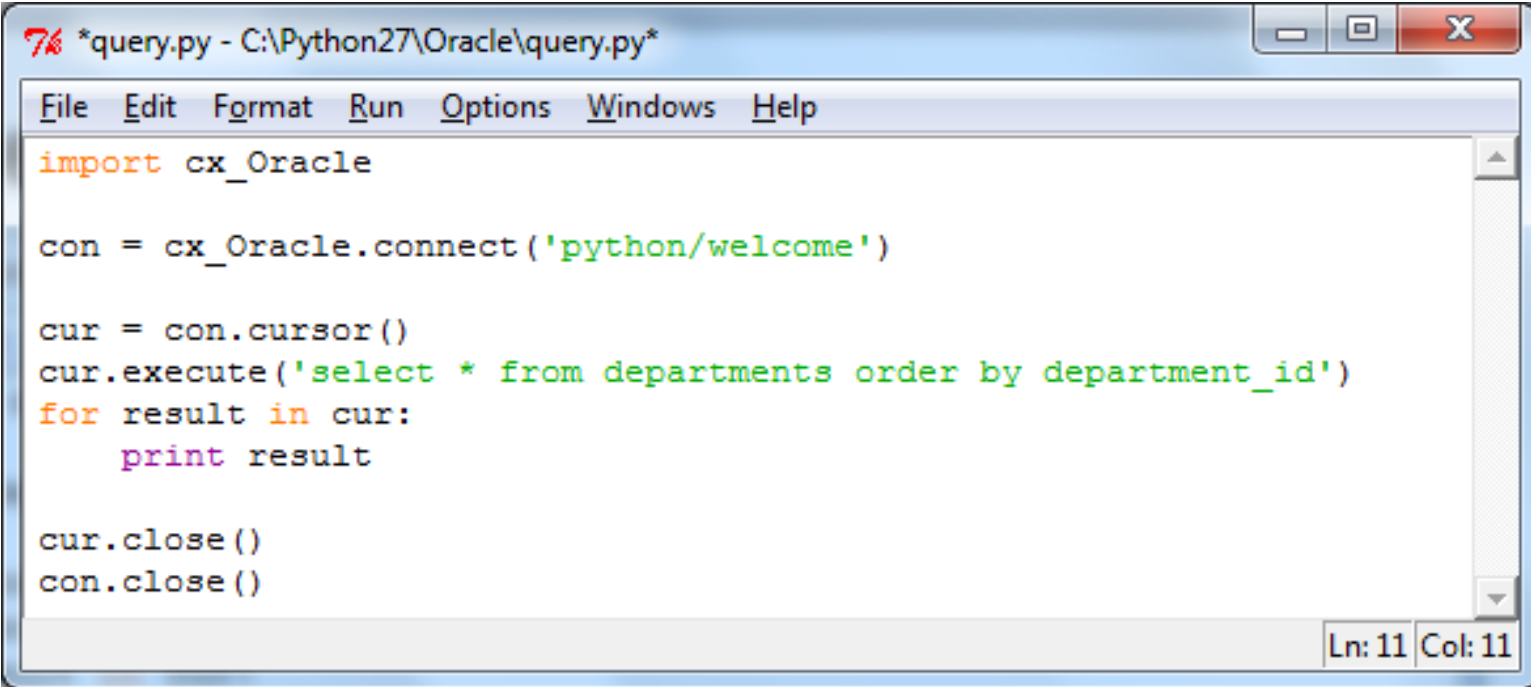
Connect Python to Oracle

- First, you import the `cx_Oracle` module to provide the API for accessing your Oracle database.
- Then, you create an object "`con`" for a specific connection, in which the username and password are passed via the `connect()` method.
- The "`con`" object has a "version" attribute, which tells you the version of Oracle you are using.
- The `close()` method is called to close the connection.

At Python Shell, enter the following:

```
>>> import cx_Oracle
>>> con=cx_Oracle.connect('python/welcome')
>>> print con.version
11.2.0.2.0
>>> con.close()
```

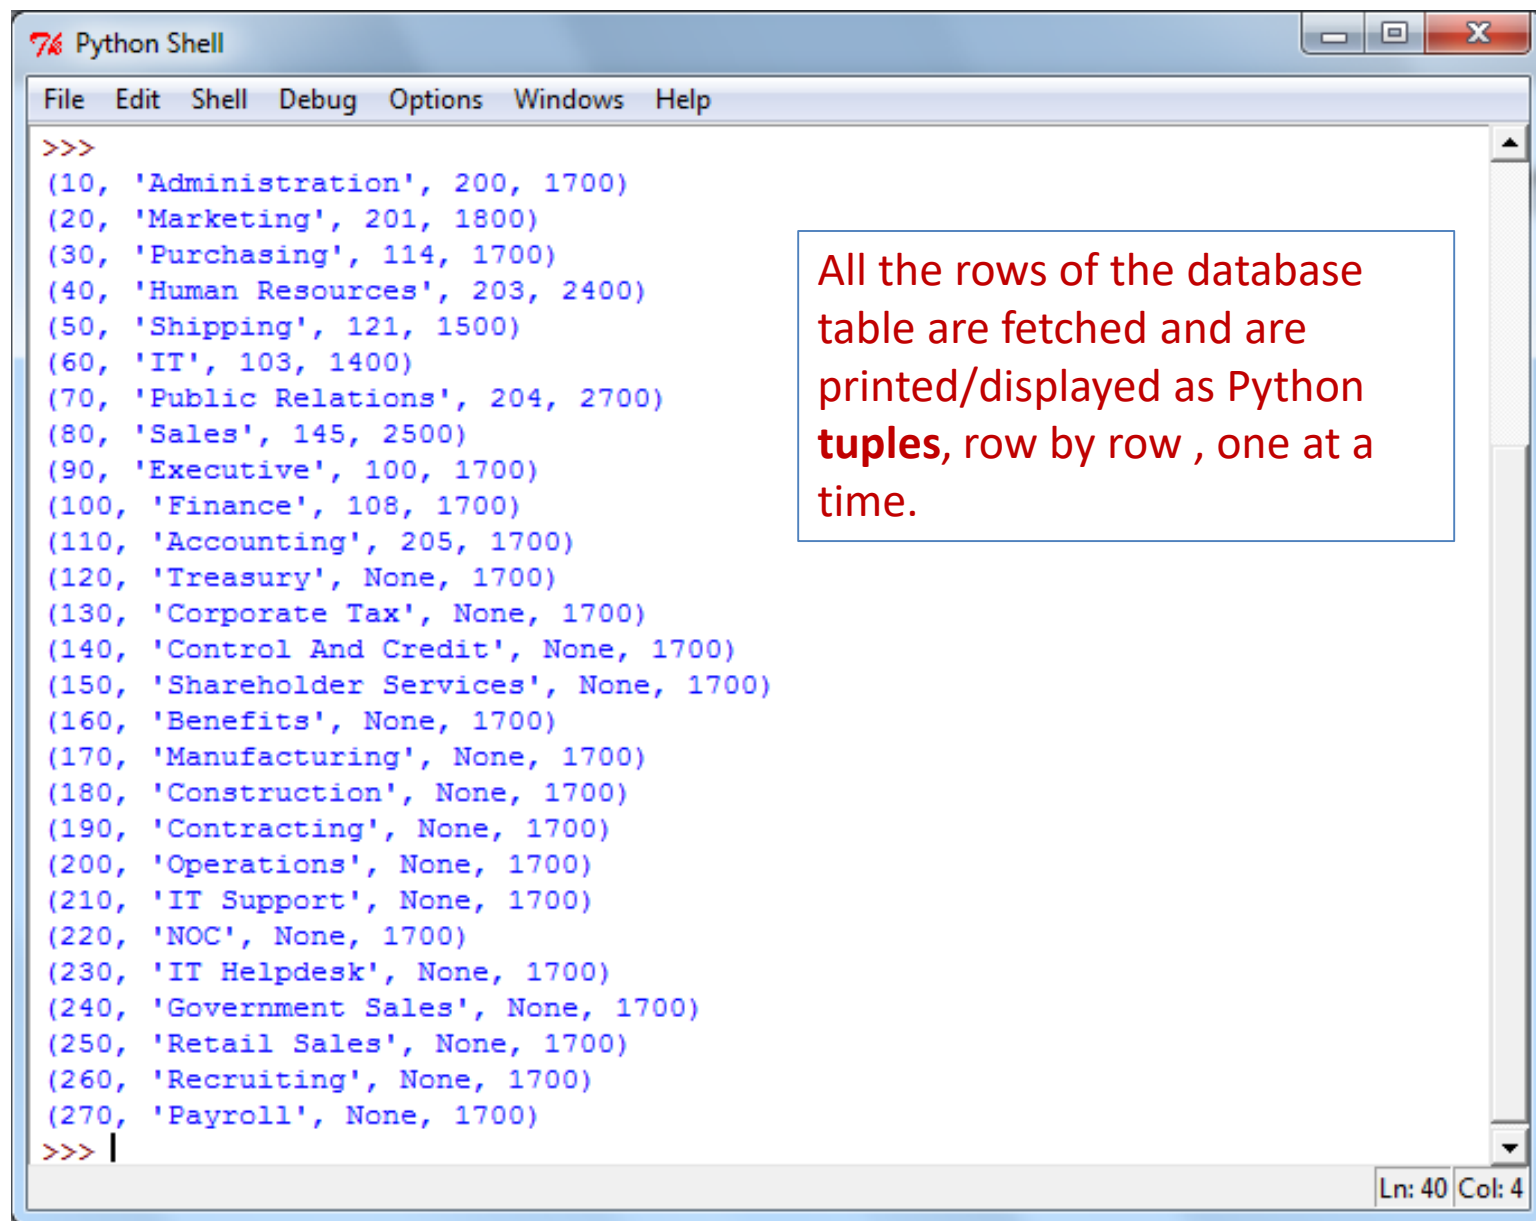
Create a Simple Query in Oracle



```
*query.py - C:\Python27\Oracle\query.py*  
File Edit Format Run Options Windows Help  
import cx_Oracle  
  
con = cx_Oracle.connect('python/welcome')  
  
cur = con.cursor()  
cur.execute('select * from departments order by department_id')  
for result in cur:  
    print result  
  
cur.close()  
con.close()  
  
Ln: 11 Col: 11
```

- The `cursor()` method opens a cursor for statements to use; a cursor is a pointer used to fetch rows from a result set.
- The `execute()` method parses and executes the SQL statement.
- The `for` loop fetches each row from the cursor and prints it.
- At the end, you need to close both the cursor and the connection.

The Result of a Simple Query in Oracle



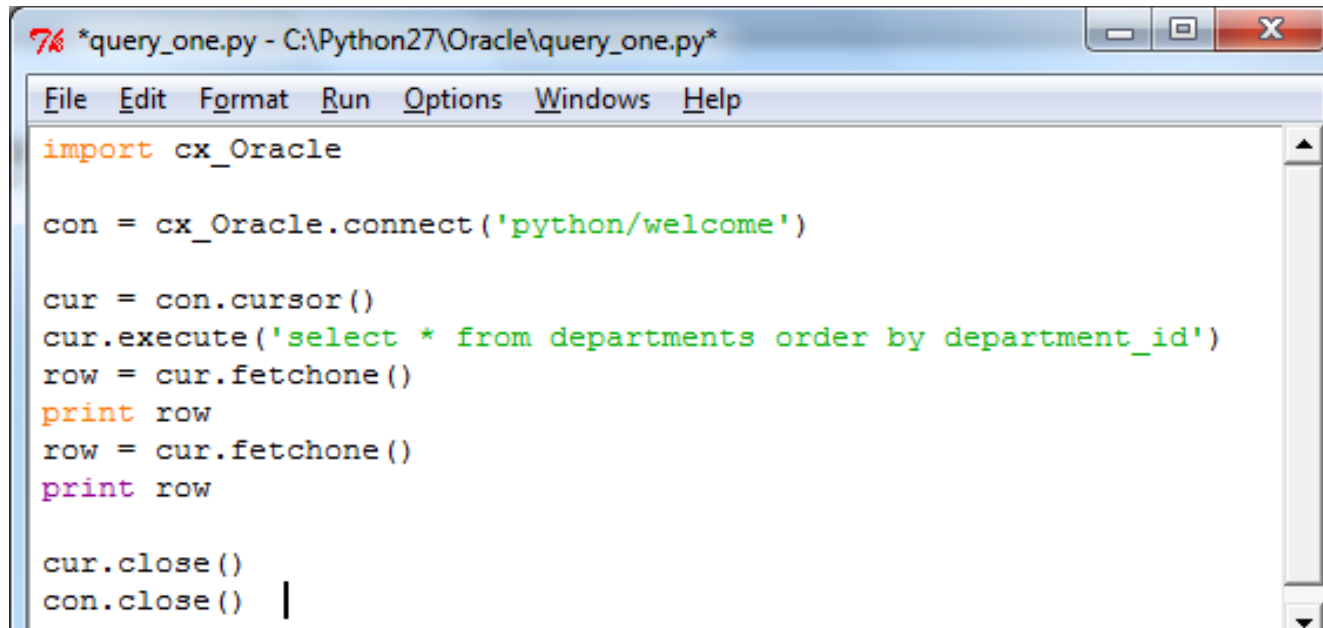
The screenshot shows a Python Shell window with a menu bar (File, Edit, Shell, Debug, Options, Windows, Help) and a text area containing the output of a query. The output consists of 27 rows, each a tuple with an ID, a department name, and a salary. The department names include Administration, Marketing, Purchasing, Human Resources, Shipping, IT, Public Relations, Sales, Executive, Finance, Accounting, Treasury, Corporate Tax, Control And Credit, Shareholder Services, Benefits, Manufacturing, Construction, Contracting, Operations, IT Support, NOC, IT Helpdesk, Government Sales, Retail Sales, Recruiting, and Payroll. The salary values are mostly 1700, with some higher values like 2400 and 2500. The status bar at the bottom right indicates 'Ln: 40 Col: 4'.

```
>>>
(10, 'Administration', 200, 1700)
(20, 'Marketing', 201, 1800)
(30, 'Purchasing', 114, 1700)
(40, 'Human Resources', 203, 2400)
(50, 'Shipping', 121, 1500)
(60, 'IT', 103, 1400)
(70, 'Public Relations', 204, 2700)
(80, 'Sales', 145, 2500)
(90, 'Executive', 100, 1700)
(100, 'Finance', 108, 1700)
(110, 'Accounting', 205, 1700)
(120, 'Treasury', None, 1700)
(130, 'Corporate Tax', None, 1700)
(140, 'Control And Credit', None, 1700)
(150, 'Shareholder Services', None, 1700)
(160, 'Benefits', None, 1700)
(170, 'Manufacturing', None, 1700)
(180, 'Construction', None, 1700)
(190, 'Contracting', None, 1700)
(200, 'Operations', None, 1700)
(210, 'IT Support', None, 1700)
(220, 'NOC', None, 1700)
(230, 'IT Helpdesk', None, 1700)
(240, 'Government Sales', None, 1700)
(250, 'Retail Sales', None, 1700)
(260, 'Recruiting', None, 1700)
(270, 'Payroll', None, 1700)
>>> |
```

Ln: 40 Col: 4

All the rows of the database table are fetched and are printed/displayed as Python **tuples**, row by row , one at a time.

Fetch Data From Database One Row a Time



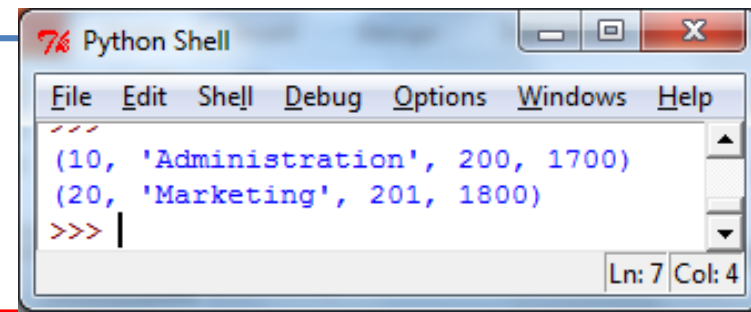
```
*query_one.py - C:\Python27\Oracle\query_one.py*
File Edit Format Run Options Windows Help
import cx_Oracle

con = cx_Oracle.connect('python/welcome')

cur = con.cursor()
cur.execute('select * from departments order by department_id')
row = cur.fetchone()
print row
row = cur.fetchone()
print row

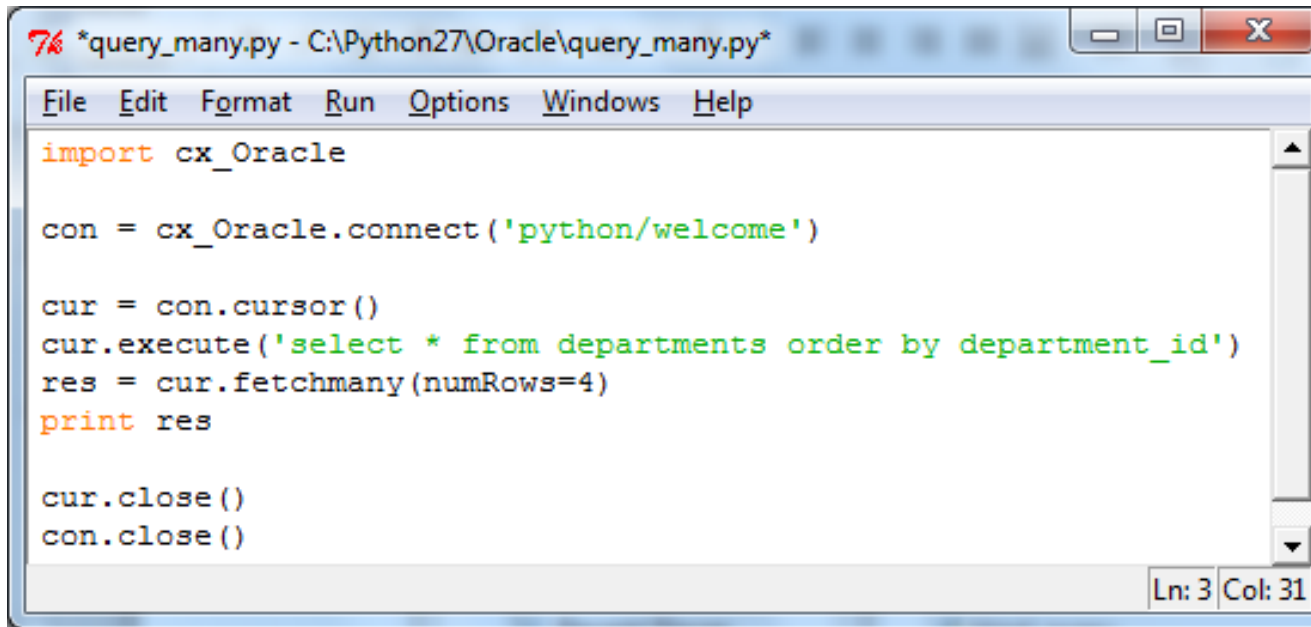
cur.close()
con.close() |
```

- The `fetchone()` method is used to return just a single row as a tuple.
- When the method is called multiple times, consecutive rows are returned.



```
Python Shell
File Edit Shell Debug Options Windows Help
(10, 'Administration', 200, 1700)
(20, 'Marketing', 201, 1800)
>>> |
Ln: 7 Col: 4
```

Fetch a Specified Number of Rows From a Database



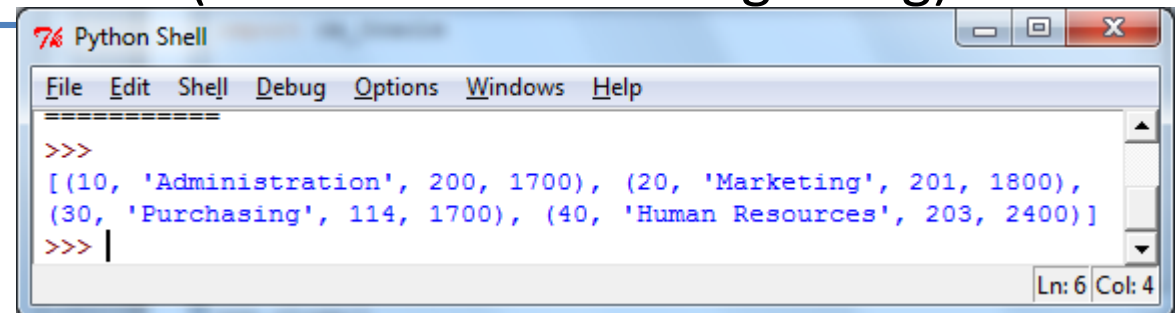
```
*query_many.py - C:\Python27\Oracle\query_many.py*
File Edit Format Run Options Windows Help
import cx_Oracle

con = cx_Oracle.connect('python/welcome')

cur = con.cursor()
cur.execute('select * from departments order by department_id')
res = cur.fetchmany(numRows=4)
print res

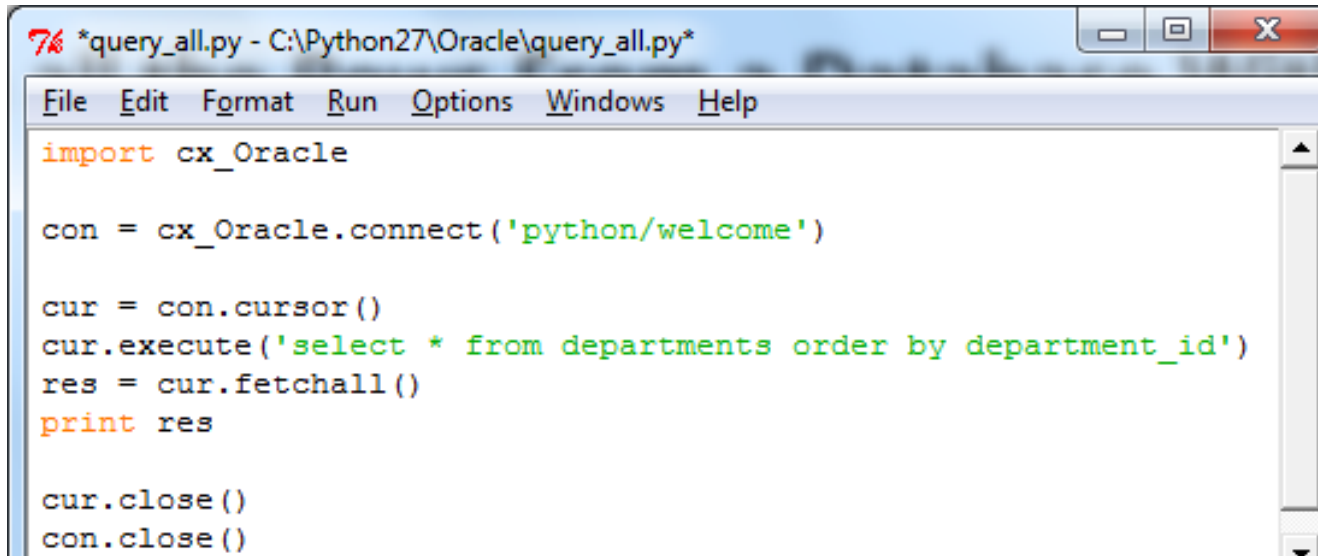
cur.close()
con.close()
Ln: 3 Col: 31
```

- The `fetchmany()` method returns a list of tuples, each of which is a row.
- The `numRows` (here=4) parameter specifies the number of rows you want to return (counted from the beginning).



```
Python Shell
File Edit Shell Debug Options Windows Help
=====
>>>
[(10, 'Administration', 200, 1700), (20, 'Marketing', 201, 1800),
(30, 'Purchasing', 114, 1700), (40, 'Human Resources', 203, 2400)]
>>> |
Ln: 6 Col: 4
```

Fetch all the Rows From a Database With fetchall()



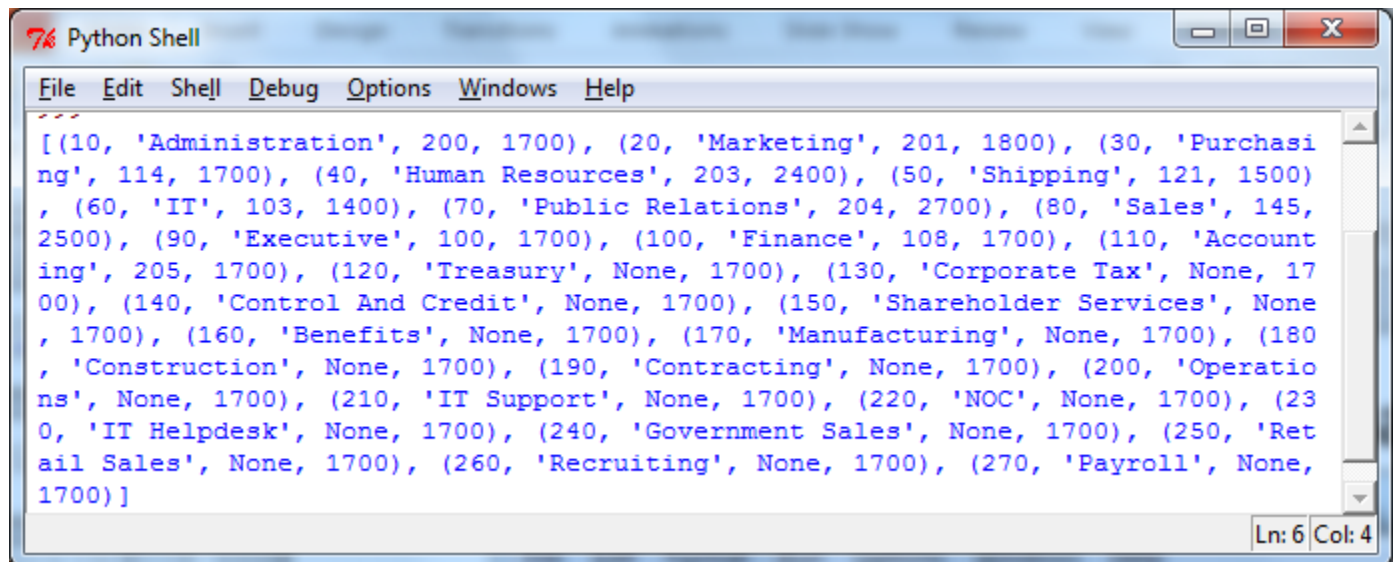
```
*query_all.py - C:\Python27\Oracle\query_all.py*
File Edit Format Run Options Windows Help
import cx_Oracle

con = cx_Oracle.connect('python/welcome')

cur = con.cursor()
cur.execute('select * from departments order by department_id')
res = cur.fetchall()
print res

cur.close()
con.close()
```

- The `fetchall()` method returns all the rows in a list of tuples.



```
Python Shell
File Edit Shell Debug Options Windows Help

[(10, 'Administration', 200, 1700), (20, 'Marketing', 201, 1800), (30, 'Purchasing', 114, 1700), (40, 'Human Resources', 203, 2400), (50, 'Shipping', 121, 1500), (60, 'IT', 103, 1400), (70, 'Public Relations', 204, 2700), (80, 'Sales', 145, 2500), (90, 'Executive', 100, 1700), (100, 'Finance', 108, 1700), (110, 'Accounting', 205, 1700), (120, 'Treasury', None, 1700), (130, 'Corporate Tax', None, 1700), (140, 'Control And Credit', None, 1700), (150, 'Shareholder Services', None, 1700), (160, 'Benefits', None, 1700), (170, 'Manufacturing', None, 1700), (180, 'Construction', None, 1700), (190, 'Contracting', None, 1700), (200, 'Operations', None, 1700), (210, 'IT Support', None, 1700), (220, 'NOC', None, 1700), (230, 'IT Helpdesk', None, 1700), (240, 'Government Sales', None, 1700), (250, 'Retail Sales', None, 1700), (260, 'Recruiting', None, 1700), (270, 'Payroll', None, 1700)]
Ln: 6 Col: 4
```

Fetch all the Rows From a Database With fetchall(), a slight modification

```

7% query_all_sep.py - C:\Python27\Oracle\query_all_sep.py
File Edit Format Run Options Windows Help
import cx_Oracle

con = cx_Oracle.connect('pythonhol/welcome')

cur = con.cursor()
cur.execute('select * from departments order by department_id')
res = cur.fetchall()
for r in res:
    print r

cur.close()
con.close()
Ln: 13 Col: 0

```

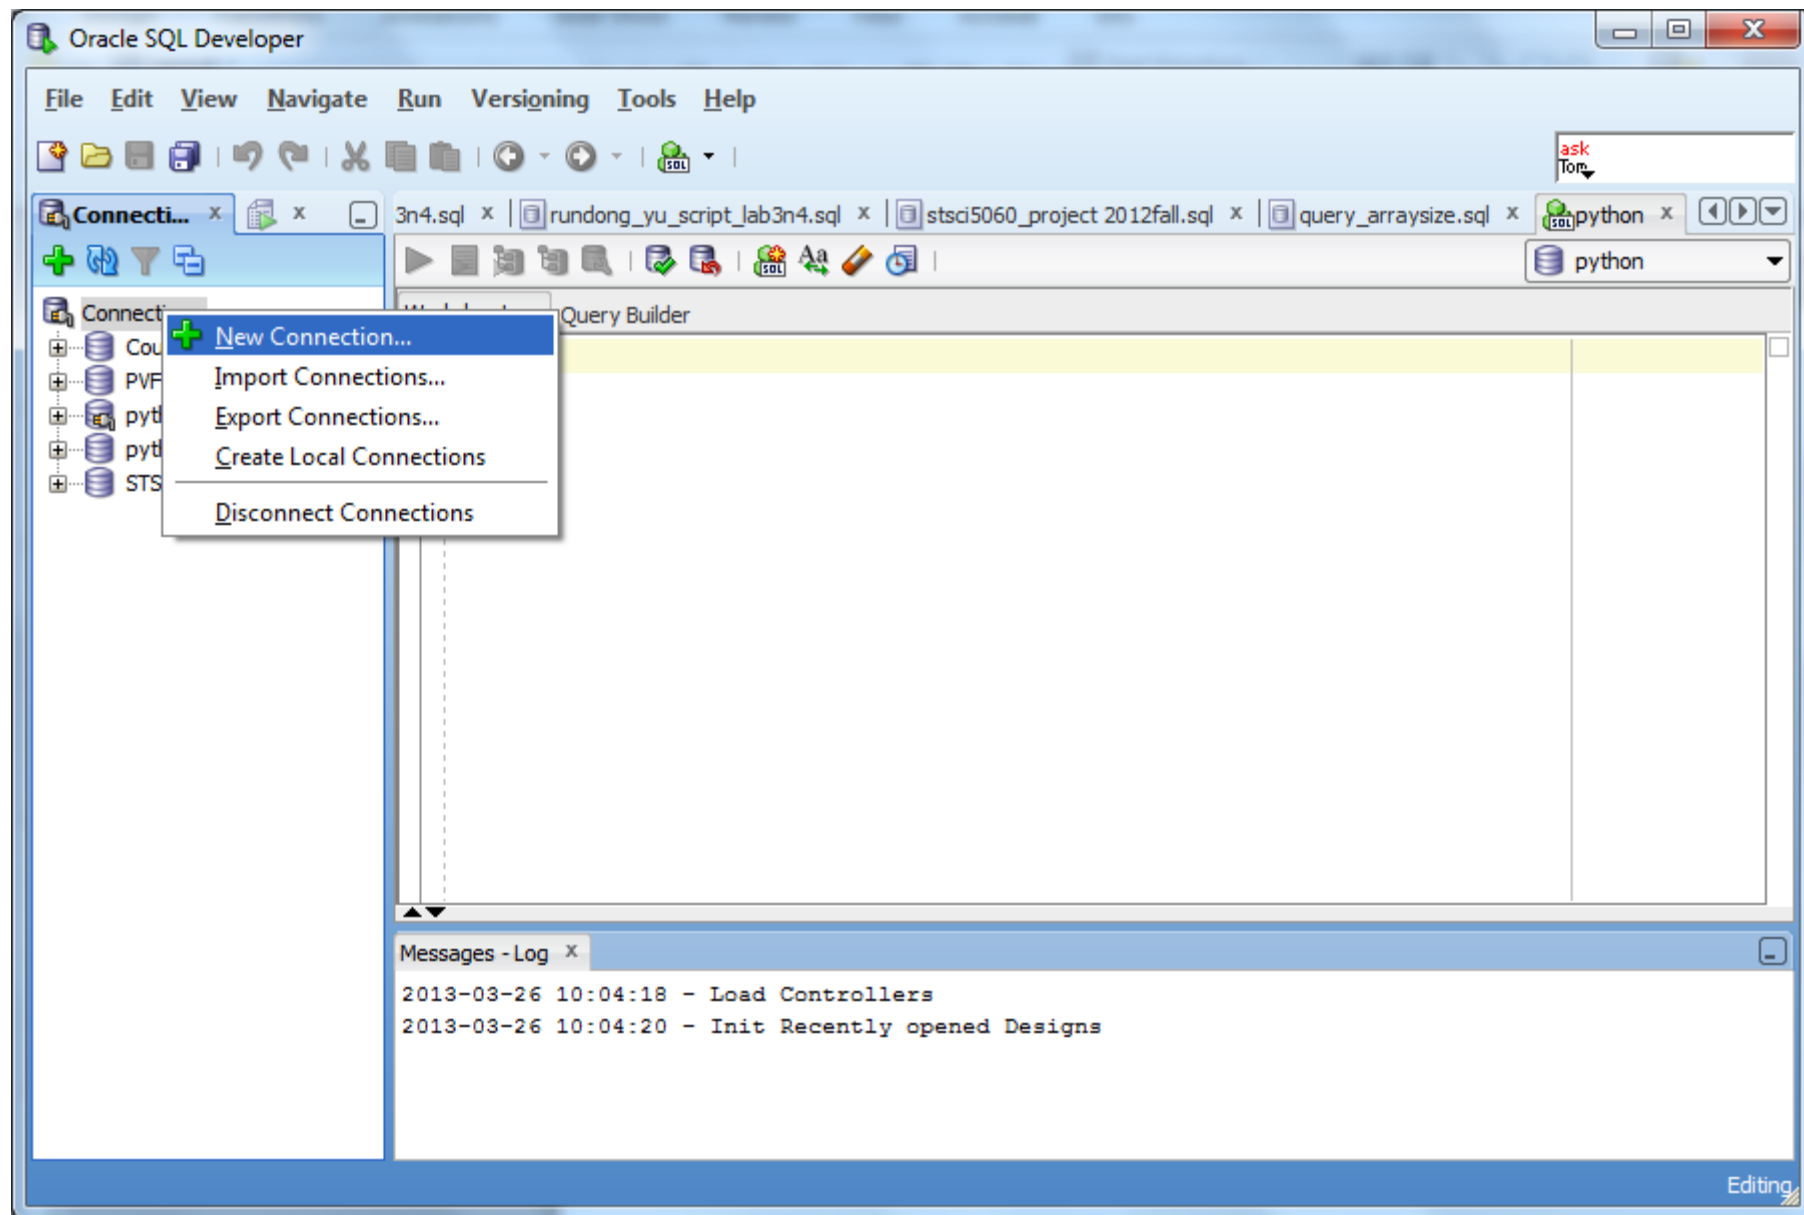
The `fetchall()` method returns all the rows in a list of tuples; each of which is now printed separately. This is the same result as displayed on Slide 11.

```

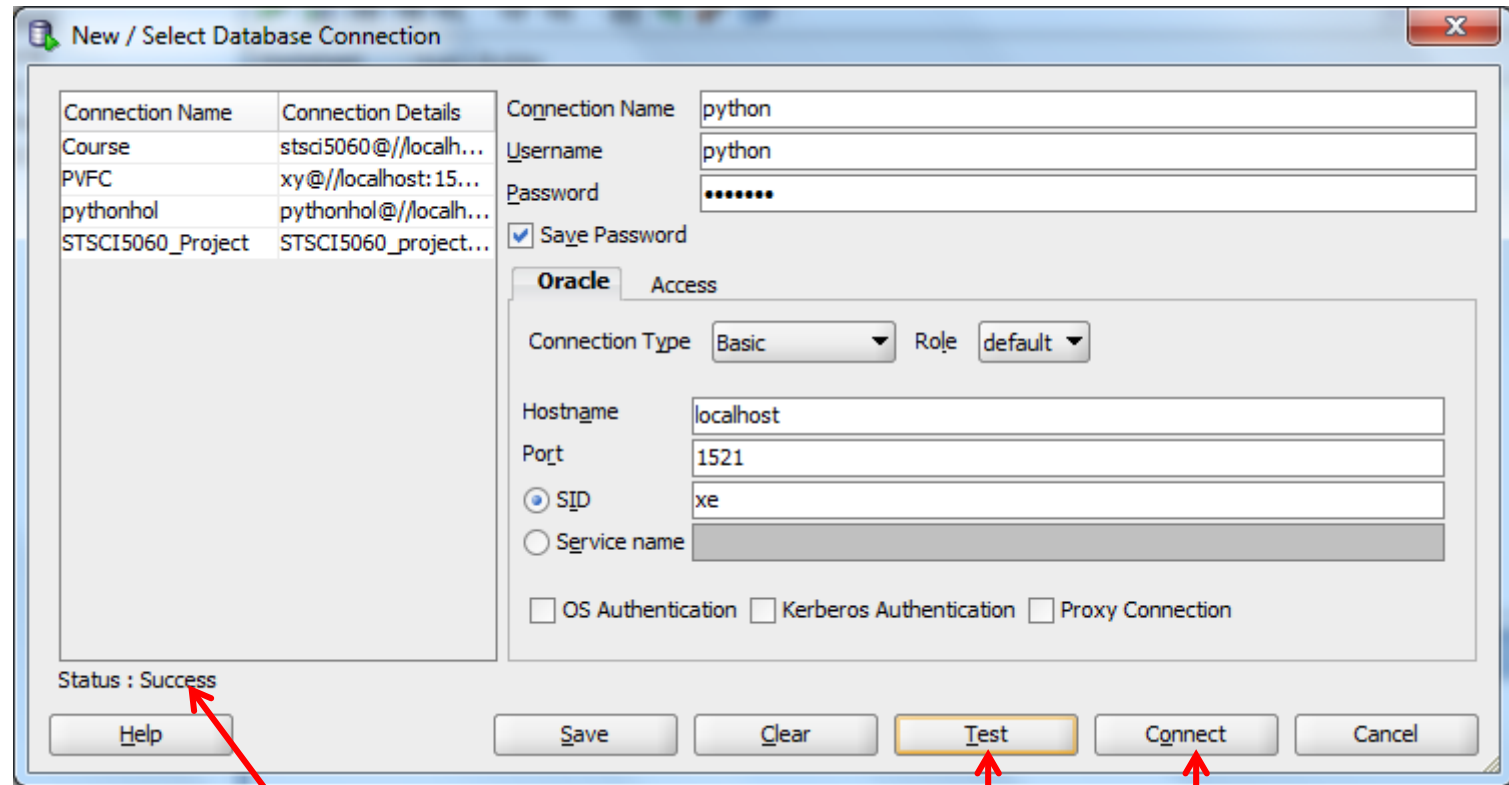
7% Python Shell
File Edit Shell Debug Options Windows Help
(10, 'Administration', 200, 1700)
(20, 'Marketing', 201, 1800)
(30, 'Purchasing', 114, 1700)
(40, 'Human Resources', 203, 2400)
(50, 'Shipping', 121, 1500)
(60, 'IT', 103, 1400)
(70, 'Public Relations', 204, 2700)
(80, 'Sales', 145, 2500)
(90, 'Executive', 100, 1700)
(100, 'Finance', 108, 1700)
(110, 'Accounting', 205, 1700)
(120, 'Treasury', None, 1700)
(130, 'Corporate Tax', None, 1700)
(140, 'Control And Credit', None, 1700)
(150, 'Shareholder Services', None, 1700)
(160, 'Benefits', None, 1700)
(170, 'Manufacturing', None, 1700)
(180, 'Construction', None, 1700)
(190, 'Contracting', None, 1700)
(200, 'Operations', None, 1700)
(210, 'IT Support', None, 1700)
(220, 'NOC', None, 1700)
(230, 'IT Helpdesk', None, 1700)
(240, 'Government Sales', None, 1700)
(250, 'Retail Sales', None, 1700)
(260, 'Recruiting', None, 1700)
(270, 'Payroll', None, 1700)
Ln: 32 Col: 4

```

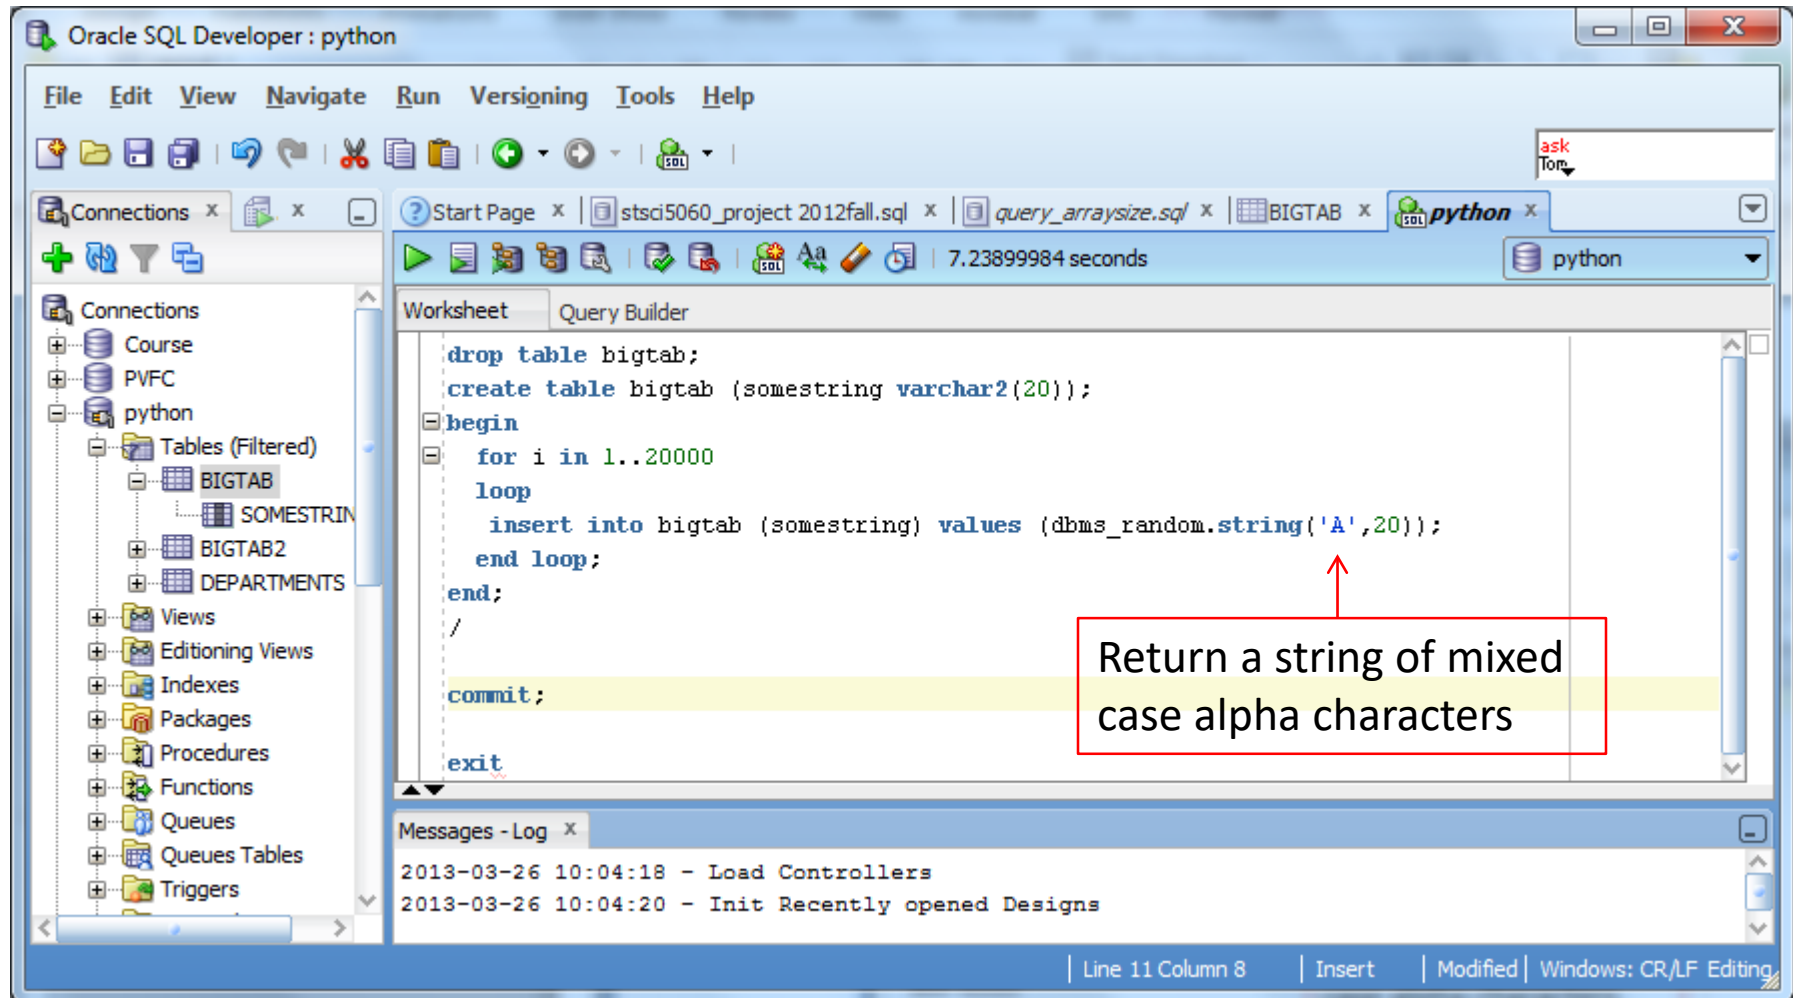
Use Oracle's SQL Developer



Use Oracle's SQL Developer (cont'd)

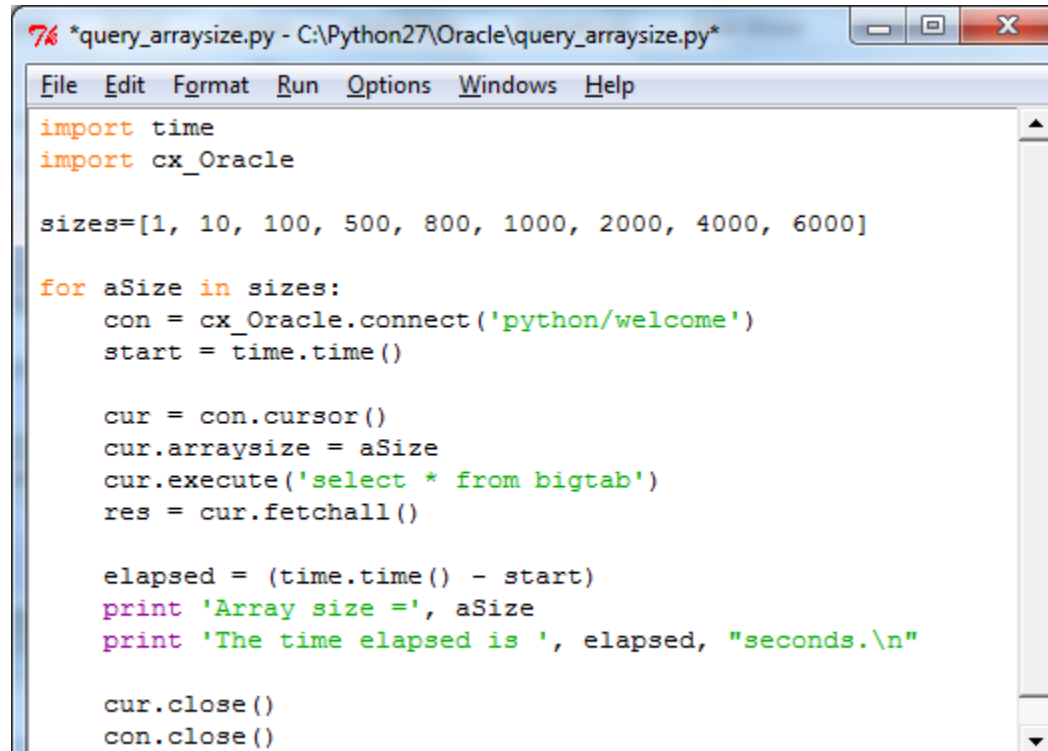


Use Oracle's SQL Developer to Create a Big Table



Improve Query Performance

Increase the number of rows returned in each batch from Oracle to the Python program can improve the query performance.



```
7% *query_arraysize.py - C:\Python27\Oracle\query_arraysize.py*
File Edit Format Run Options Windows Help
import time
import cx_Oracle

sizes=[1, 10, 100, 500, 800, 1000, 2000, 4000, 6000]

for aSize in sizes:
    con = cx_Oracle.connect('python/welcome')
    start = time.time()

    cur = con.cursor()
    cur.arraysize = aSize
    cur.execute('select * from bigtab')
    res = cur.fetchall()

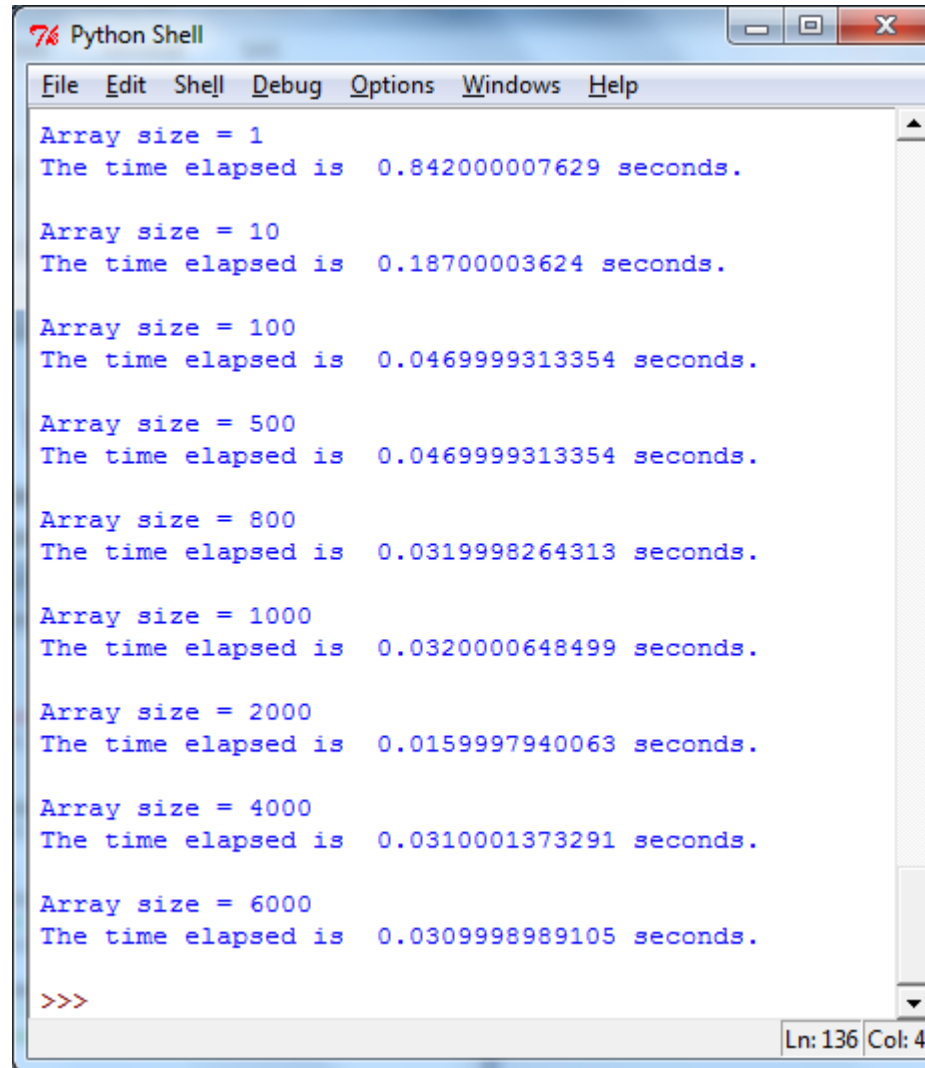
    elapsed = (time.time() - start)
    print 'Array size =', aSize
    print 'The time elapsed is ', elapsed, "seconds.\n"

    cur.close()
    con.close()
```

- The 'time' module is used to measure elapsed time of the query.
- The arraysize is set to different values in **sizes** list. This causes batches of **arraysize** records at a time to be returned from the Oracle DB to a **cache** in Python, which reduces the number of "**roundtrips**" made to the DB, often reducing network load and reducing the number of context switches on the database server.
- The fetchall() method reads from the cache before requesting more data from Oracle.

Improve Query Performance (cont'd)

The output for comparing the performance



A screenshot of a Python Shell window titled "Python Shell". The window has a menu bar with "File", "Edit", "Shell", "Debug", "Options", "Windows", and "Help". The main text area displays the following output:

```
Array size = 1  
The time elapsed is 0.842000007629 seconds.  
  
Array size = 10  
The time elapsed is 0.18700003624 seconds.  
  
Array size = 100  
The time elapsed is 0.0469999313354 seconds.  
  
Array size = 500  
The time elapsed is 0.0469999313354 seconds.  
  
Array size = 800  
The time elapsed is 0.0319998264313 seconds.  
  
Array size = 1000  
The time elapsed is 0.0320000648499 seconds.  
  
Array size = 2000  
The time elapsed is 0.0159997940063 seconds.  
  
Array size = 4000  
The time elapsed is 0.0310001373291 seconds.  
  
Array size = 6000  
The time elapsed is 0.0309998989105 seconds.  
  
>>>
```

The status bar at the bottom right shows "Ln: 136 Col: 4".

Using Bind Variables

Bind variables enable you to re-execute statements with new values, without the overhead of **reparsing** the statement. As a result, it improves code reusability.

```
7% *bind_query.py - D:/STSCI4060/STSCI4060_SP2013/Lectures/File7/bind_query.py*
File Edit Format Run Options Windows Help
import cx_Oracle

con = cx_Oracle.connect('python/welcome')

cur = con.cursor()
cur.prepare('select * from departments where department_id = :id')

deptList = [10,20,30,210,110,180,260]
for deptid in deptList:
    cur.execute(None, {'id': deptid})
    res = cur.fetchone()
    print res

cur.close()
con.close()
```

- Here the bind variable is **:id**. The statement is only prepared once but executed seven times with different values for the WHERE clause.
- The special symbol '**None**' is used in place of the statement text argument to execute() because the prepare() method has already set the statement. The second argument to the execute() call is a Python dictionary.
- In the execute() calls, this dictionary has the values of the deptList.

```
7% Python Shell
File Edit Shell Debug Options Windows Help
>>>
(10, 'Administration', 200, 1700)
(20, 'Marketing', 201, 1800)
(30, 'Purchasing', 114, 1700)
(210, 'IT Support', None, 1700)
(110, 'Accounting', 205, 1700)
(180, 'Construction', None, 1700)
(260, 'Recruiting', None, 1700)
>>> |
```

Create a Table in Oracle Within Python

```

createEmptab.py - D:/STSCI4060/STSCI4060_SP2013/Lectures/File7/createEmptab.py
File Edit Format Run Options Windows Help
import cx_Oracle

con = cx_Oracle.connect('python/welcome')
cur=con.cursor()
cur.execute('create table emptab (id number, salary number, age number)')
cur.close()
con.close()
Ln: 8 Col: 0

```

Oracle SQL Developer : Table PYTHON.EMPTAB@python

File Edit View Navigate Run Versioning Tools Help

Connections x | Start Page x | stsci5060_project 2012fall.sql x | EMPTAB x | python~2 x

Columns Data Constraints Grants Statistics Triggers Flashback Dependencies Details Partitions Indexes

Columns

COLUMN_NAME	DATA_TYPE	NULLABLE	DATA_DEFAULT	COLUMN_ID	COMMENTS
ID	NUMBER	Yes	(null)	1 (null)	
SALARY	NUMBER	Yes	(null)	2 (null)	
AGE	NUMBER	Yes	(null)	3 (null)	

Messages - Log x

2013-03-26 10:04:18 - Load Controllers
2013-03-26 10:04:20 - Init Recently opened Designs

python | PYTHON | EMPTAB Editing

Populate the Oracle Table Within Python

```
import cx_Oracle

con = cx_Oracle.connect('python/welcome')

rows = [ (1, 90500, 45),
          (2, 45500, 30),
          (3, 32000, 28),
          (4, 65000, 40),
          (5, 68000, 41),
          (6, 80000, 39),
          (7, 120000, 51),
          (8, 85900, 44),
          (9, 100800, 48),
          (10, 166000, 60)]

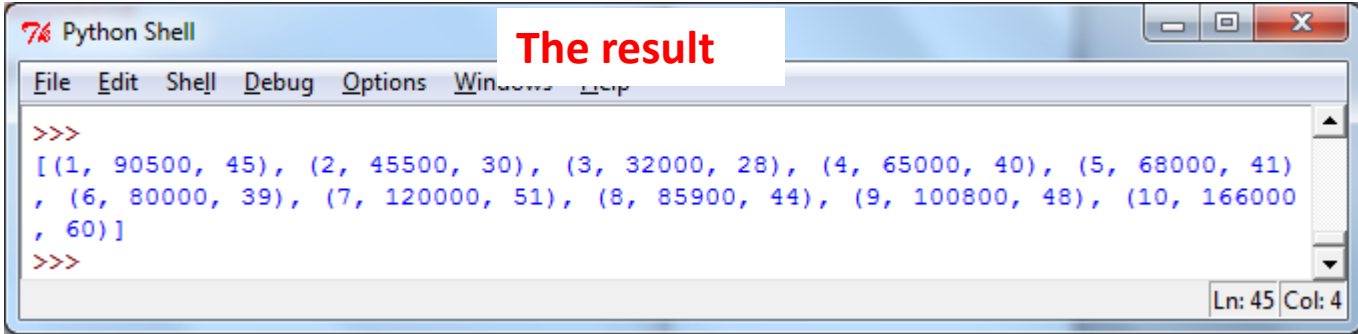
cur = con.cursor()
cur.bindarraysize = 10
cur.setinputsizes(int, 8, 3)
cur.executemany("insert into emptab(id, salary, age) values (:c1, :c2, :c3)", rows)

# Now query the results back
cur2 = con.cursor()
cur2.execute('select * from emptab')
res = cur2.fetchall()
print res

cur.close()
cur2.close()
con.close()
```

Populate the Oracle Table Within Python (cont'd)

You got a result in Python shell



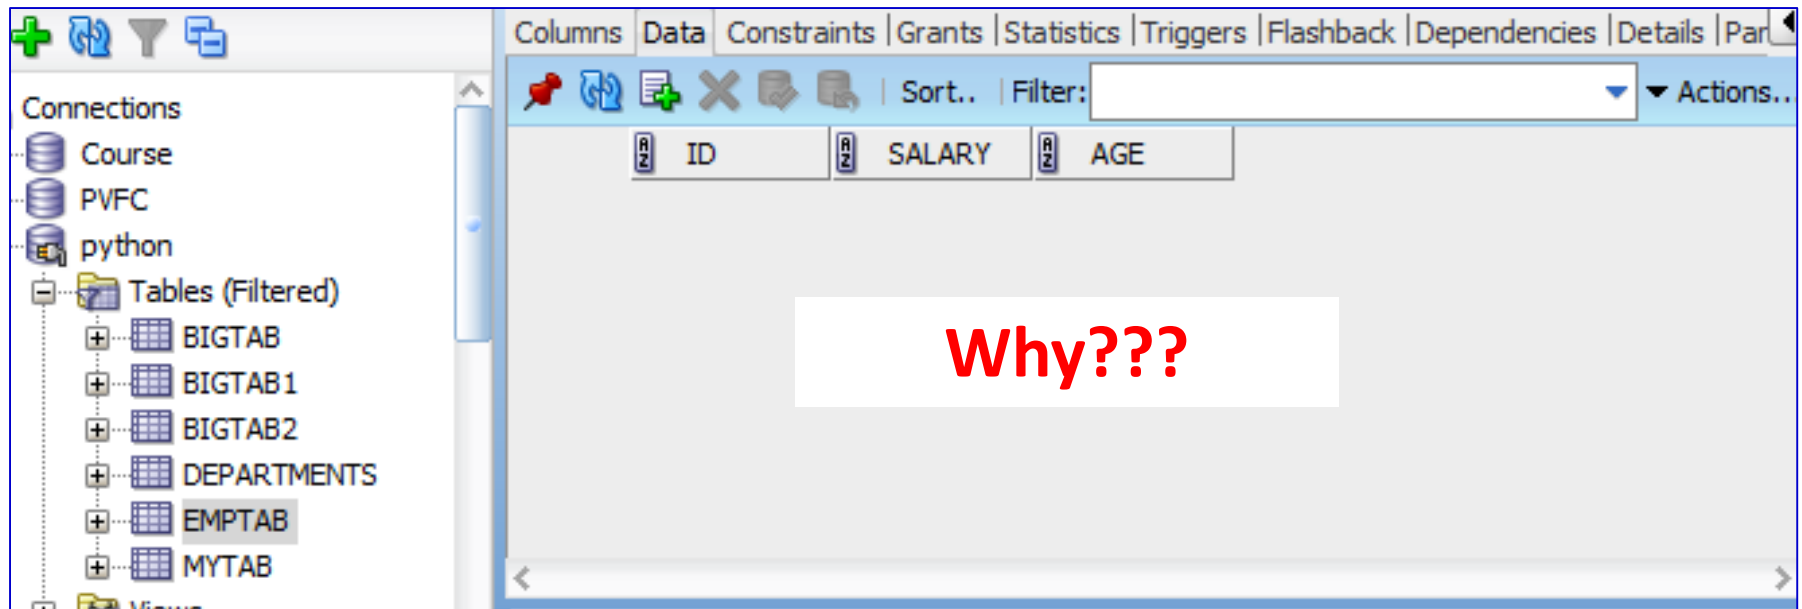
The screenshot shows a Python Shell window with a menu bar (File, Edit, Shell, Debug, Options, Windows, Help) and a text area containing the following Python code and its output:

```
>>> [(1, 90500, 45), (2, 45500, 30), (3, 32000, 28), (4, 65000, 40), (5, 68000, 41),  
, (6, 80000, 39), (7, 120000, 51), (8, 85900, 44), (9, 100800, 48), (10, 166000  
, 60)]  
>>>
```

A red box labeled "The result" points to the output list. The status bar at the bottom right indicates "Ln: 45 Col: 4".

Populate the Oracle Table Within Python (cont'd)

But you did not see the data in the
Oracle table EMPTAB



Populate the Oracle Table Within Python (cont'd)

Because you did not commit the changes.

When you manipulate data in an Oracle database (insert, etc.), the changed or new data are only available within your database session until it is committed to the database by running the **commit()** method/command. When the changes are committed to the database, they are then available to other users and sessions. This is a database transaction. Note that the commit() method is used on the connection, not on the cursor.

Populate the Oracle Table Within Python (cont'd)

Committing the changes – creating an transaction

```
import cx_Oracle

con = cx_Oracle.connect('python/welcome')

rows = [ (1, 90500, 45),
          (2, 45500, 30),
          (3, 32000, 28),
          (4, 65000, 40),
          (5, 68000, 41),
          (6, 80000, 39),
          (7, 120000, 51),
          (8, 85900, 44),
          (9, 100800, 48),
          (10, 166000, 60)]

cur = con.cursor()
cur.bindarraysize = 10
cur.setinputsizes(int, 8, 3)
cur.executemany("insert into emptab(id, salary, age) values (:c1, :c2, :c3)", rows)

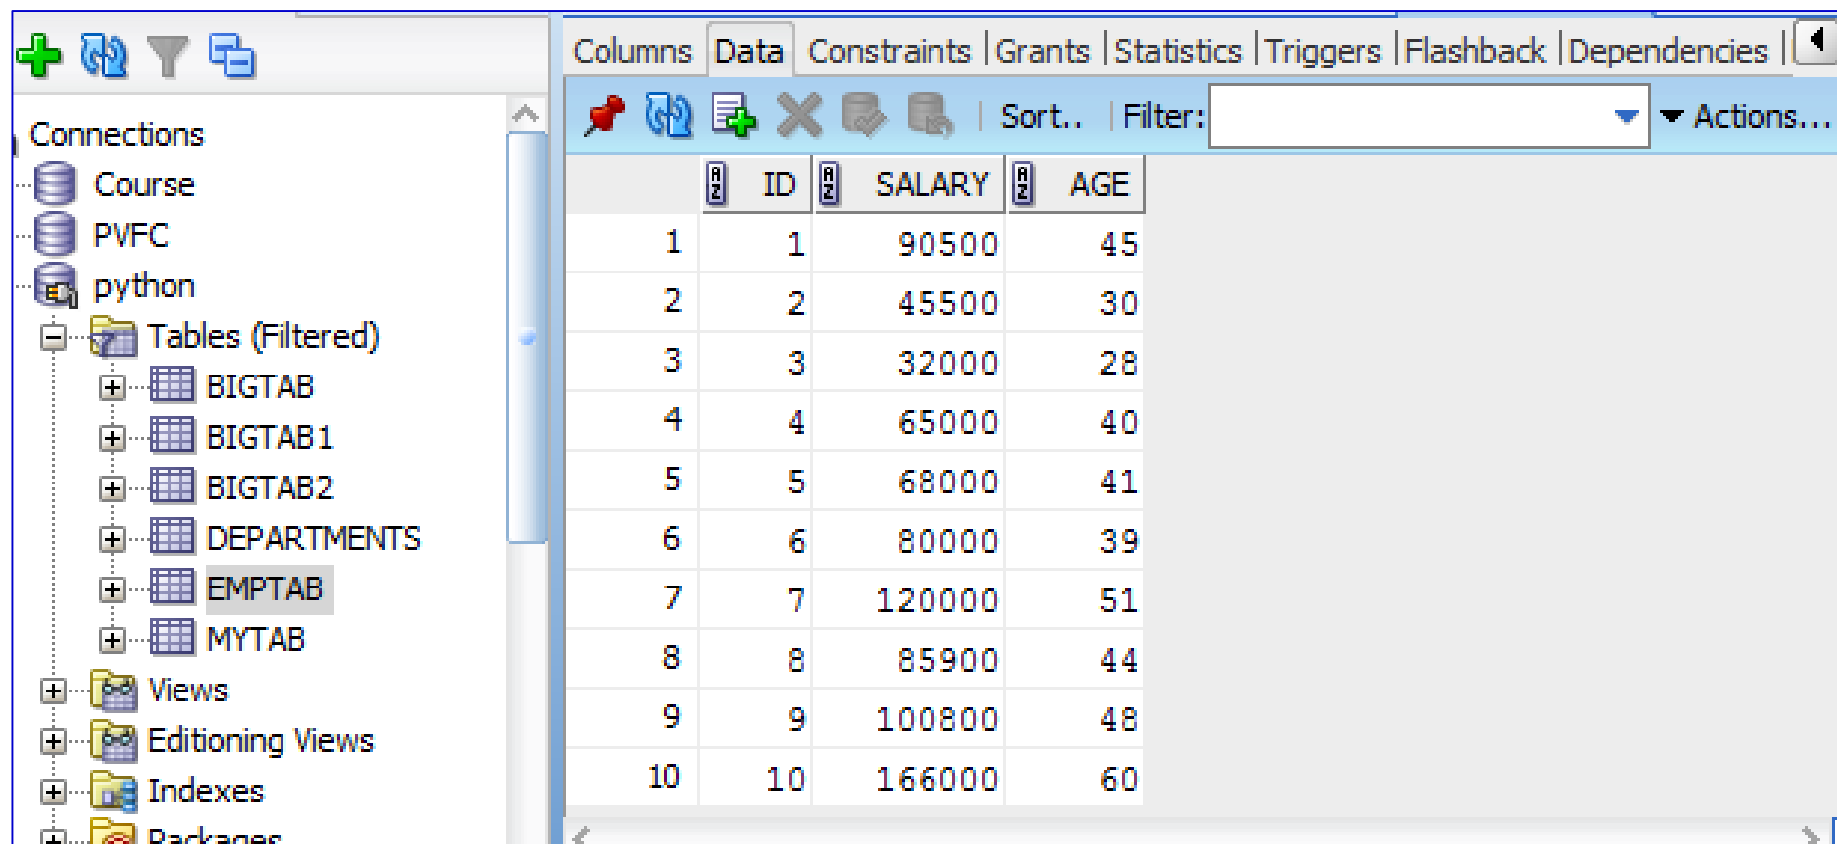
#Commit the changes to finish the DB transaction
con.commit()

# Now query the results back from the Oracle DB
cur2 = con.cursor()
cur2.execute('select * from emptab')
res = cur2.fetchall()
print res

cur.close()
cur2.close()
con.close()
```

Populate the Oracle Table Within Python (cont'd)

Now you can see the data in the Oracle table EMPTAB



The screenshot shows the Oracle SQL Developer interface. On the left, the 'Connections' tree is expanded to show the 'python' connection, and the 'Tables (Filtered)' folder is selected, displaying a list of tables including EMPTAB. The main pane shows the 'Data' tab for the EMPTAB table, displaying a grid of 10 rows of data. The columns are ID, SALARY, and AGE. The data is as follows:

	ID	SALARY	AGE
1	1	90500	45
2	2	45500	30
3	3	32000	28
4	4	65000	40
5	5	68000	41
6	6	80000	39
7	7	120000	51
8	8	85900	44
9	9	100800	48
10	10	166000	60

Update an Oracle Table Through Python

Task: Raise the salary of the employee with id=7 from 120000 to 155000 due to a promotion in position.

```
import cx_Oracle

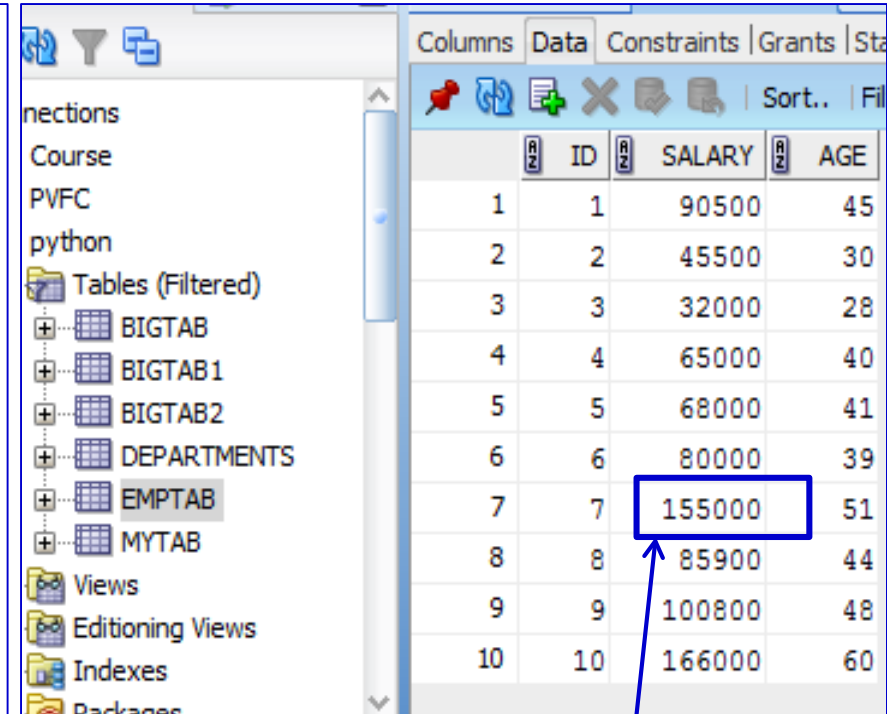
con = cx_Oracle.connect('python/welcome')

cur = con.cursor()
cur.execute('''update emptab
              set salary=155000
              where id=7''')

#Commit the changes to finish the DB transaction
con.commit()

# Now query the results back from the Oracle DB
cur2 = con.cursor()
cur2.execute('select * from emptab')
res = cur2.fetchall()
print res

cur.close()
cur2.close()
con.close()
```



The screenshot shows the Oracle SQL Developer interface. On the left, the 'Connections' pane shows a connection named 'python'. Below it, the 'Tables (Filtered)' list includes BIGTAB, BIGTAB1, BIGTAB2, DEPARTMENTS, EMPTAB (selected), and MYTAB. On the right, the 'Data' tab displays the contents of the 'emptab' table. The table has columns ID, SALARY, and AGE. The row for ID 7 shows a salary of 155000, which is highlighted with a blue box. A blue arrow points from this box to a text box at the bottom right that says 'The update made'.

	ID	SALARY	AGE
1	1	90500	45
2	2	45500	30
3	3	32000	28
4	4	65000	40
5	5	68000	41
6	6	80000	39
7	7	155000	51
8	8	85900	44
9	9	100800	48
10	10	166000	60

Note: use triple quotation marks to keep the Oracle syntax format.

The update made

Obtain Data From the Oracle DB

Retrieve the **salary** data from the EMPTAB table in the Oracle DB and do some simple statistics of the data.

```
import cx_Oracle
import scipy

con = cx_Oracle.connect('python/welcome')

cur = con.cursor()
cur.execute('select salary from emptab')
sal=cur.fetchall()
num=range(len(sal))
print 'The salary list is:'
for i in num:
    print '$' + str(sal[i][0]),
print '\nThe average salary is $' + str(scipy.mean(sal)) + '.'
print 'The maximum salary is $' + str(max(sal)[0]) + '.'
print 'The minimum salary is $' + str(min(sal)[0]) + '.'

cur.close()
con.close()
```

```
The salary list is:
$90500 $45500 $32000 $65000 $68000 $80000 $155000 $85900 $100800 $166000
The average salary is $88870.0.
The maximum salary is $166000.
The minimum salary is $32000.
```

Obtain Data From the Oracle DB, Another Example

Retrieve the **salary** and **age** data from the EMPTAB table in the Oracle DB and do a regression analysis by importing the Python program, mylr.py, which we wrote earlier in the class (with a small modification).

```
import cx_Oracle
import mylr

con = cx_Oracle.connect('python/welcome')

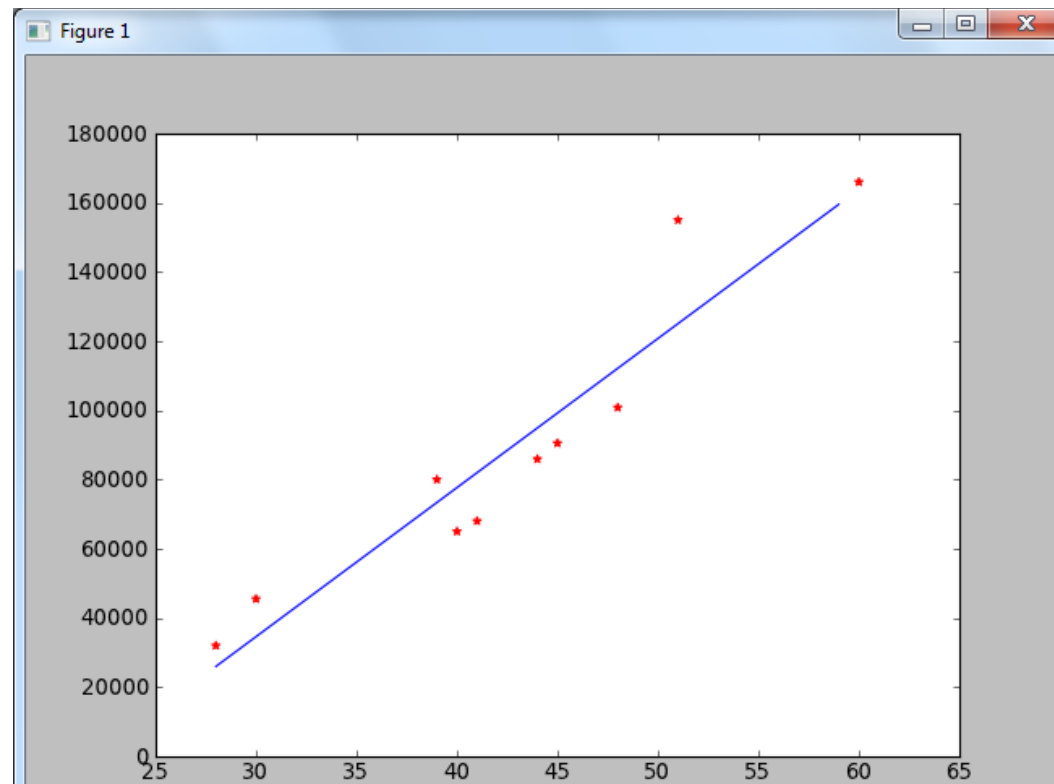
cur = con.cursor()
cur.execute('select salary, age  from emptab')
data=cur.fetchall()
num=range(len(data))
a=[]
s=[]
for i in num:
    s.append(data[i][0])
    a.append(data[i][1])

print 'The salary list (y variable): ', s
print 'The age list (x variable): ', a

print '\nThe linear regression result:'
mylr.lmr(a,s)

cur.close()
con.close()
```

Obtain Data From the Oracle DB, Another Example



```
*Python Shell*
File Edit Shell Debug Options Windows Help
The salary list (y variable): [90500, 45500, 32000, 65000, 68000, 80000, 155000,
85900, 100800, 166000]
The age list (x variable): [45, 30, 28, 40, 41, 39, 51, 44, 48, 60]

The linear regression result:
The number of observations = 10
Slope = 4308
Intercept = -94651
The equation is: y = -94651 + 4308 x
```