

Lab 7 - Random Effects and Least Significant Differences

Lab Goals

1. Derive the distribution of conditional normal random vectors; and simulate these.
2. Provide optimal predictions of random effects in a Randomized Block design via a conditional argument.
3. Implement random effects models in R using the `lme4` package.
4. Explore blocked and longitudinal designs.

Conditional Normal Random Vectors

Here we will examine the distribution of a normal random vector, when you know part of the vector. That is, we we have

$$\begin{pmatrix} y_1 \\ y_2 \end{pmatrix} \sim N \left(\begin{pmatrix} \mu_1 \\ \mu_2 \end{pmatrix}, \begin{pmatrix} \Sigma_{11} & \Sigma_{12} \\ \Sigma_{21} & \Sigma_{22} \end{pmatrix} \right)$$

and we see y_1 , what does that tell us about y_2 ?

This can be done directly by looking at the joint density of y_1 and y_2 and treating y_1 as fixed, but there is a simpler argument looking at $z = y_2 - Ay_1$ for a specific choice

$$A = \Sigma_{21}\Sigma_{11}^{-1}$$

1. Show that $cov(z, y_1) = 0$ and hence that z is independent of y_1 . What is $E(z)$?
2. By expressing $y_2 = z + Ay_1$ and using the independence above, find $E(y_2|y_1)$.
3. Taking the same approach show that $var(y_2|y_1) = var(z)$.
4. Find $var(z)$
5. The code below conducts a simulation study. First, we generate 1000 bivariate normal random variables, with variance 2 and covariance 1:

```
eta = rnorm(1000)  # This is a common effect to both y's
y1 = rnorm(1000) + eta
y2 = rnorm(1000)+eta
```


Now we'll look at taking the y_2 that correspond to y_1 close to the value 1. We can illustrate this with the following plot

```
plot(y1,y2)
abline(v = c(0.5,1.5))

#Take only those y2 near y1=1

cy2 = y2[ (y1>0.5) & (y1<1.5) ]

# Plot these in red
points( y1[ (y1>0.5) & (y1<1.5) ], cy2, col='red')
```



According to our conditional variance formula, these should be Normal with mean $1/2$ and variance $(2 - 1/2) = 3/2$

```
mean(cy2)
```

```
## [1] 0.3954537
```

```
var(cy2)
```


```
## [1] 1.391946
```

Of course, these will be a bit distorted because we used y_1 's that were only close to 1. If we look at the histogram, in comparison to the expected density.

```
hist(cy2,20,prob=TRUE)
```

```
xpts = seq(-4,4,len=1001)
```

```
lines(xpts, dnorm(xpts,mean=0.5,sd=sqrt(3/2)))
```



Best Linear Unbiased Predictors in a Randomized Block Design

Here we will use the result in a Randomized Blocks Design in which we have

$$y_{ij} = \mu + \alpha_i + \beta_j + e_{ij}$$

with $e_{ij} \sim N(0, \sigma_e^2)$ and $\alpha_i \sim N(0, \sigma_a^2)$.

1. We consider estimating α_i by $\hat{\alpha}_i = \bar{y}_{i\cdot} - \mu$. Give an expression for the error, $\hat{\alpha}_i - \alpha_i$ in terms of μ , α , β and e .
2. Hence find the expected squared error $E(\hat{\alpha}_i - \alpha_i)^2$.
3. The BLUP estimates for α_i is

$$\tilde{\alpha}_i = \frac{\sigma_a^2}{\sigma_a^2 + \sigma_e^2/r} (\bar{y}_{i\cdot} - \mu)$$

Derive the expected squared error $E(\tilde{\alpha}_i - \alpha_i)^2$

4. Show that the error for $\tilde{\alpha}_i$ is smaller than for $\hat{\alpha}_i$.
5. *bonus* We have assumed we know μ in the calculations above. What if we replace it by $\bar{y}_{..}$?

A Randomized Blocks Analysis

In Lab 6, we analyzed the distance that 4 different brands of golfballs travelled when hit. In fact, there were 10 different golfers in this study, given by the first column. We'll first load in these data

```
golf = read.table('golfballs.dat')
names(golf) = c('golfer', 'brand', 'dist')
golf$golfer = as.factor(golf$golfer)
```

1. We can express the mean-model framework for this design by kronecker products:

```
# Golfer design matrix
Xg = diag(10)%x%matrix(1,4,1)

# Check this
Xg[1:12,1:4]
```

```
##      [,1] [,2] [,3] [,4]
## [1,]    1    0    0    0
## [2,]    1    0    0    0
## [3,]    1    0    0    0
## [4,]    1    0    0    0
## [5,]    0    1    0    0
## [6,]    0    1    0    0
## [7,]    0    1    0    0
## [8,]    0    1    0    0
## [9,]    0    0    1    0
## [10,]   0    0    1    0
## [11,]   0    0    1    0
## [12,]   0    0    1    0
```

```
# Design for golf ball brands
Xb = matrix(1,10,1)%x%diag(4)

# Check
Xb[1:12,]
```

```
##      [,1] [,2] [,3] [,4]
## [1,]    1    0    0    0
## [2,]    0    1    0    0
## [3,]    0    0    1    0
## [4,]    0    0    0    1
## [5,]    1    0    0    0
## [6,]    0    1    0    0
## [7,]    0    0    1    0
## [8,]    0    0    0    1
## [9,]    1    0    0    0
## [10,]   0    1    0    0
## [11,]   0    0    1    0
## [12,]   0    0    0    1
```

To be able to estimate a model, we have to use a reference category for each factor, and then add an intercept.

```
X = cbind( rep(1,40), Xg[, -1], Xb[, -1])
```

2. To get parameters for each golfer and each brand, we can simply do linear regression

```
betahat = solve( t(X)%*%X)%*%t(X)%*%golf$dist
```

We can compare this to a simple linear regression model

```
mod1 = lm(dist~golfer+brand,data=golf)
mod1$coefficients
```

```
## (Intercept)      golfer2      golfer3      golfer4      golfer5      golfer6
##      203.7025      39.5750      15.5500      28.7000      -2.7750      44.9750
##      golfer7      golfer8      golfer9      golfer10      brandB      brandC
##       6.6000      35.4500      19.1250      46.2750       6.1300      18.2800
##      brandD
##      -6.3200
```

```
t(betahat)
```

```
##           [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10]
## [1,] 203.7025 39.575 15.55 28.7 -2.775 44.975  6.6 35.45 19.125 46.275
##           [,11] [,12] [,13]
## [1,]  6.13 18.28 -6.32
```

3. We can also express this in terms of averages for each group. To do this, we need to

a. Extract the over-all mean

```
ydd = mean(golf$dist)
ydd
```

```
## [1] 231.5725
```

b. Average for each golfer

```
yid = t(Xg)%*%golf$dist/4
t(yid)
```

```
##           [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9]
## [1,] 208.225 247.8 223.775 236.925 205.45 253.2 214.825 243.675 227.35
##           [,10]
## [1,] 254.5
```

c. Average each brand

```
ydj = t(Xb)%*%golf$dist/10
t(ydj)
```

```
##           [,1] [,2] [,3] [,4]
## [1,] 227.05 233.18 245.33 220.73
```

Notice that if we predict the first element of the data set by

```
yid[1] + ydj[1] - ydd
```

```
## [1] 203.7025
```

We get the same thing as

```
mod1$fit[1]
```

```
##      1
## 203.7025
```

bonus How do you map the coefficients of the linear model onto these averages? (Note that this isn't completely straightforward).

d. We can now estimate variances

```
# For errors
sig2e = sum( (golf$dist - Xg%%yid - Xb%%yjd + ydd)^2)/(3*9)

# For golfers
sig2g = sum( (yid - ydd)^2 )/9 - sig2e/4
```


e. And look at BLUPS for golfers

```
gblup = sig2g/(sig2g+sig2e/4)*(yid-ydd)
t(gblup)
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## [1,] -22.99516 15.98261 -7.679828  5.271725 -25.72829 21.30112 -16.49476
##           [,8]      [,9]     [,10]
## [1,] 11.91986 -4.158778 22.5815
```

We can plot these, too

```
plot(yid-ydd,gblup)
abline(c(0,1))
```



Here the shrinkage is very slight because between-golfer variance is much larger than the error variance.

4. To perform a mixed effects analysis in R, we use the package `lme4`

```
library(lme4)
```

```
## Warning: package 'lme4' was built under R version 3.5.1
```

```
## Loading required package: Matrix
```

```
mod2 = lmer(dist~brand + (1|golfer),data=golf)
summary(mod2)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula: dist ~ brand + (1 | golfer)
## Data: golf
##
## REML criterion at convergence: 257.4
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -2.16064 -0.38059  0.00973  0.52941  1.42909
##
## Random effects:
## Groups Name Variance Std.Dev.
## golfer (Intercept) 330.32  18.175
## Residual          20.25   4.499
## Number of obs: 40, groups: golfer, 10
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)  227.050      5.921  38.347
## brandB         6.130      2.012   3.046
## brandC        18.280      2.012   9.084
## brandD        -6.320      2.012  -3.141
##
## Correlation of Fixed Effects:
##      (Intr) brandB brandC
## brandB -0.170
## brandC -0.170  0.500
## brandD -0.170  0.500  0.500
```

Notice that we estimate variances between golfers and hits within golfers as we obtained manually.

We can also get the BLUPs for each golfer

```
ranef(mod2)
```

```
## $golfer
## (Intercept)
## 1 -22.995163
## 2  15.982611
## 3  -7.679828
## 4   5.271725
## 5 -25.728285
## 6  21.301119
## 7 -16.494764
## 8  11.919861
## 9  -4.158778
## 10 22.581501
```

These are not the least squares estimates, but they are the same as the blups we calculated.

A Simulation Exercise

If you get time. Simulate a one-way random effects model with 5 levels and 2 observations per level and set $\mu = 0$, $\sigma_a^2 = \sigma_e^2 = 1$.

For example

```
Xa = diag(5)%x%matrix(1,2,1)
levs = as.factor( (1:5)%x%rep(1,2) )

# Random effects
alpha = rnorm(5)


# Create data
y = Xa*%alpha + rnorm(10)

# Let's run a model
mod3 = lmer(y~1+(1|levs))

# See what this looks like
summary(mod3)

## Linear mixed model fit by REML ['lmerMod']
## Formula: y ~ 1 + (1 | levs)
##
## REML criterion at convergence: 27.8
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.4149 -0.7716 -0.2887  0.7717  1.5587
##
## Random effects:
##   Groups   Name      Variance Std.Dev.
##   levs      (Intercept) 4.312e-16 2.077e-08
##   Residual                9.926e-01 9.963e-01
## Number of obs: 10, groups:  levs, 5
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)  0.02037    0.31506   0.065

# And check estimation of random effects
plot(alpha,ranef(mod3)$levs[[1]])
```

Simulate this 1000 times. Show that the estimate of μ that we would find is more variable than we would have if we modeled this with fixed effects. Show that the blups for this model have smaller squared error (over different α s) than using the un-shrunken estimates.

A Longitudinal Data Analysis

First, let's look at an interaction between categorical and continuous covariates.

1. Suppose that X^S is a matrix coding the indicator for different levels of a (fixed) factor S . So that X_1^S is an indicator that $S = 1$ and so forth. We will also take t to be a continuous covariate – say, the time at which a measurement is taken.

Writing

$$y = X^S \beta_x + t X^S \beta_t + e$$

Show that this gives each level of S its own slope and intercept.

2. In a fixed effects setting in which X^S is given with reference coding, write down a model that extracts a reference linear regression and interpret the interaction effects.
3. For random effects, we don't use reference level coding since each level of S is random. Specify X and Z in the model

$$y = X\beta + Zb + e$$

with $e \sim N(0, \sigma_e^2)$ and $b \sim N(0, G)$ so that each level of S has its own random slope and intercept.

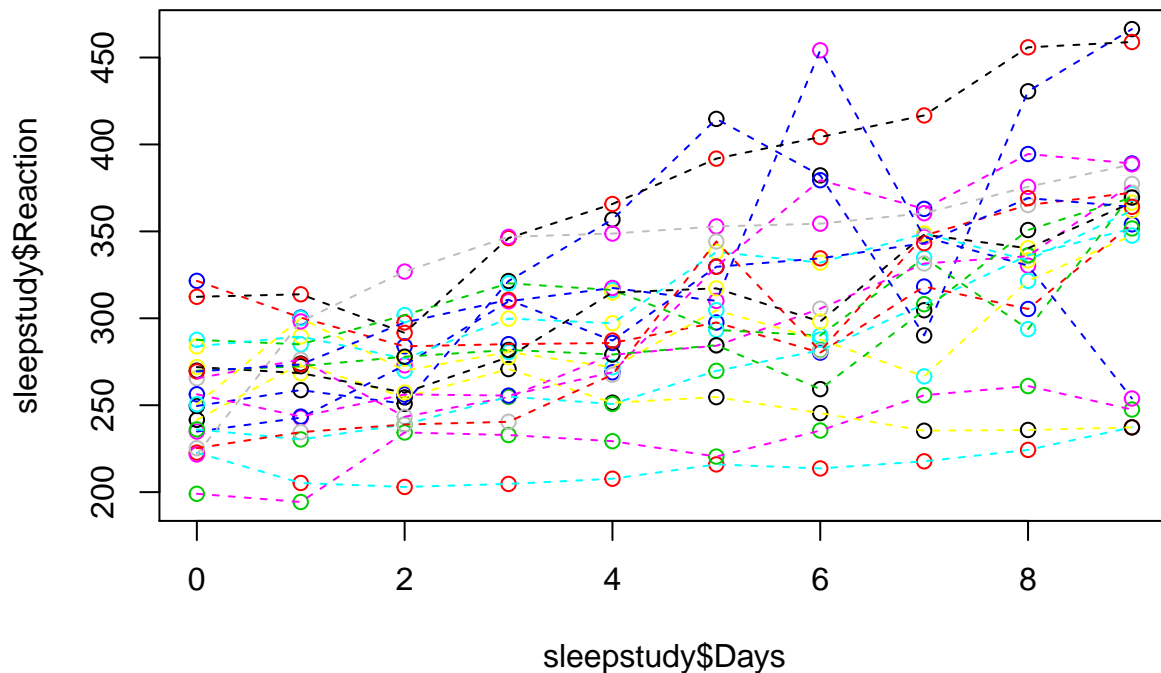
4. This model applies to data on reaction times in the data `sleepstudy` from the package `lme4`.

```
library(lme4)
data(sleepstudy)
```

Here we have 18 subjects coded in `Subject`, they were each restricted to 3 hours sleep per night for 10 nights and their reaction time (`Reaction`) measured on each of 10 days (`Day`). Here we assume that each subject's reaction time increases with longer sleep deprivation (ie, each has their own linear regression) but that this relationship is random between subjects.

a. We can see this more clearly by plotting the data. Connecting the values for each subject with lines

```
plot(sleepstudy$Days,sleepstudy$Reaction,col=sleepstudy$Subject)
for(j in levels(sleepstudy$Subject)){
  which.obs = (sleepstudy$Subject == j)  # These rows correspond to subject j
  lines(sleepstudy[which.obs,c(2,1)],col=j,lty=2)  # Plot day against time for these rows
}
```



b. This model can be fit in `lme4` as follows

```
longmod = lmer(Reaction~Days+(Days|Subject),data=sleepstudy)
summary(longmod)
```


```
## Linear mixed model fit by REML ['lmerMod']
## Formula: Reaction ~ Days + (Days | Subject)
## Data: sleepstudy
##
## REML criterion at convergence: 1743.6
##
## Scaled residuals:
##    Min       1Q   Median       3Q      Max
```

```
## -3.9536 -0.4634 0.0231 0.4634 5.1793
##
## Random effects:
## Groups   Name            Variance Std.Dev. Corr
## Subject  (Intercept) 612.09   24.740
##          Days        35.07    5.922  0.07
## Residual                654.94   25.592
## Number of obs: 180, groups: Subject, 18
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)  251.405     6.825   36.838
## Days         10.467     1.546    6.771
##
## Correlation of Fixed Effects:
##      (Intr)
## Days -0.138
```

Notice that specifying (Days|Subject) provides both an intercept and a slope for each subject.

c. We can extract and plot the random effects with

```
rand.effects = ranef(longmod)
plot(rand.effects$Subject)
```



d. By the default, the model allows there to be a correlation between each subjects slope and intercept. The (very small) positive correlation estimated here says that a subject with a longer starting reaction time is also affected more by sleep deprivation.

That correlation is pretty small, though. We can tell lme4 not to allow it with the following

```
longmod2 = lmer(Reaction~Days + (1|Subject)+(0+Days|Subject),data=sleepstudy)
summary(longmod2)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula: Reaction ~ Days + (1 | Subject) + (0 + Days | Subject)
## Data: sleepstudy
##
## REML criterion at convergence: 1743.7
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.9626 -0.4625  0.0204  0.4653  5.1860
##
## Random effects:
## Groups      Name                Variance Std.Dev.
## Subject    (Intercept)    627.57     25.051
## Subject.1 Days              35.86      5.988
## Residual                    653.58     25.565
## Number of obs: 180, groups: Subject, 18
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)  251.405      6.885   36.513
## Days         10.467      1.560    6.712
##
## Correlation of Fixed Effects:
##      (Intr)
## Days -0.184
```

Here we have to specify (0+Days|Subject) so that the second random term doesn't include an intercept. You can see that we no longer have an estimated correlation.

- e. (Beyond course material). Are these models different? We can compare their likelihoods in a formal test (see BTRY/STSCI 4090) using

```
anova(longmod,longmod2)
```

```
## refitting model(s) with ML (instead of REML)
## Data: sleepstudy
## Models:
## longmod2: Reaction ~ Days + (1 | Subject) + (0 + Days | Subject)
## longmod: Reaction ~ Days + (Days | Subject)
##      Df    AIC    BIC logLik deviance Chisq Chi Df Pr(>Chisq)
## longmod2  5 1762.0 1778.0 -876.00   1752.0
## longmod   6 1763.9 1783.1 -875.97   1751.9 0.0639      1    0.8004
```

5. In the fully study, there was a group that got 6 hours sleep instead of 3. How would you set up a model to test whether the *average* slope in the second group was lower than that in the first?