

MDBM-Chapter 7

Advanced SQL

Main Knowledge Points

- Multiple table SQL queries
- Different types of joins
- Noncorrelated and correlated subqueries
- Referential integrity in SQL
- Triggers and stored procedures, etc.

Processing Multiple Tables: JOINS

- **JOIN:** a relational operation that causes two or more tables with a common domain to be combined into a single result table or view. The common columns in joined tables are often the primary key of the dominant table and the foreign key of the dependent table in 1:M relationships.
- **JOIN implicitly:** using a WHERE clause, which is the traditional way and was the only means before SQL 1992.
- **JOIN explicitly:** using the **JOIN...ON** commands. There are a number of join operations available by combining with different key words: **INNER, OUTER, FULL, LEFT, RIGHT, CROSS** and **UNION**. Not all DBMSes support all the key words.

EQUI-JOIN

EQUI-JOIN: a join in which the joining condition is based on **equality** between values in the common columns; common columns can appear redundantly in the result table. It is normally established by the key words **INNER JOIN ... ON**

```
SELECT d.department_name, e.employee_name  
FROM departments_t d  
    INNER JOIN employees_t e  
        ON d.department_id = e.department_id  
WHERE d.department_id >= 30  
ORDER BY d.department_name;
```

OUTER JOIN

(A join in which rows that do not have matching values in common columns can also be included in the result table.)

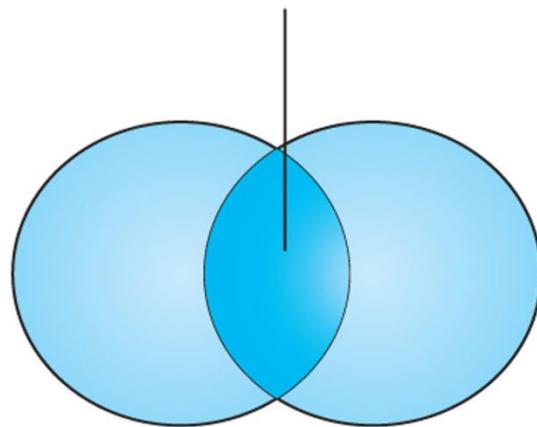
- **LEFT OUTER JOIN:** a join in which rows from the table on the left side of the “LEFT OUTER JOIN” keywords that do not have matching values in common columns are also included in the result table; otherwise, the result will be the same as the INNER JOIN.
- **RIGHT OUTER JOIN:** a join in which rows from the table on the right side of the “RIGHT OUTER JOIN” keywords that do not have matching values in common columns are also included in the result table; otherwise, the result will be the same as the INNER JOIN.
- **FULL OUTER JOIN:** a join in which rows from the tables on the both sides of the “FULL OUTER JOIN” keywords that do not have matching values in common columns are also included in the result table; otherwise, the result will be the same as the LEFT OUTER JOIN or RIGHT OUTER JOIN.

UNION JOIN

- ❑ A join that includes all the data from each table in the join and an instance for each row of each table.
- ❑ The result table contains all the columns from each table; NULLs will be assigned to the columns where there are no values.
- ❑ This command is not available in Oracle, but it is easy to achieve the same result in Oracle.

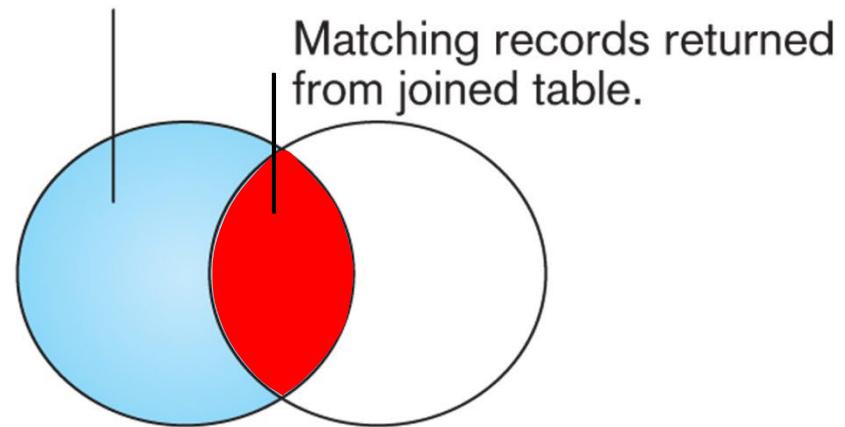
Visualization of Different Join Types with Results Returned in Shaded Area

Darker area is result returned.



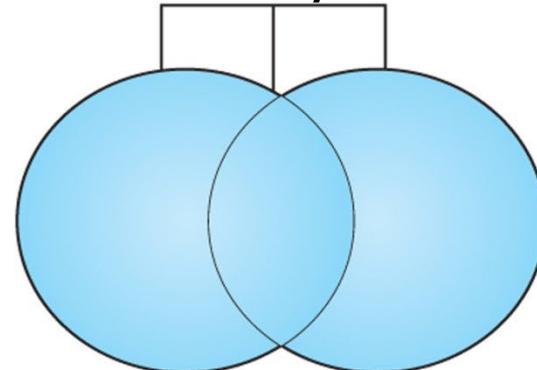
Natural Join

All the records of the left table may be returned



Left Outer Join

All the records may be returned



Union Join

Set Operations

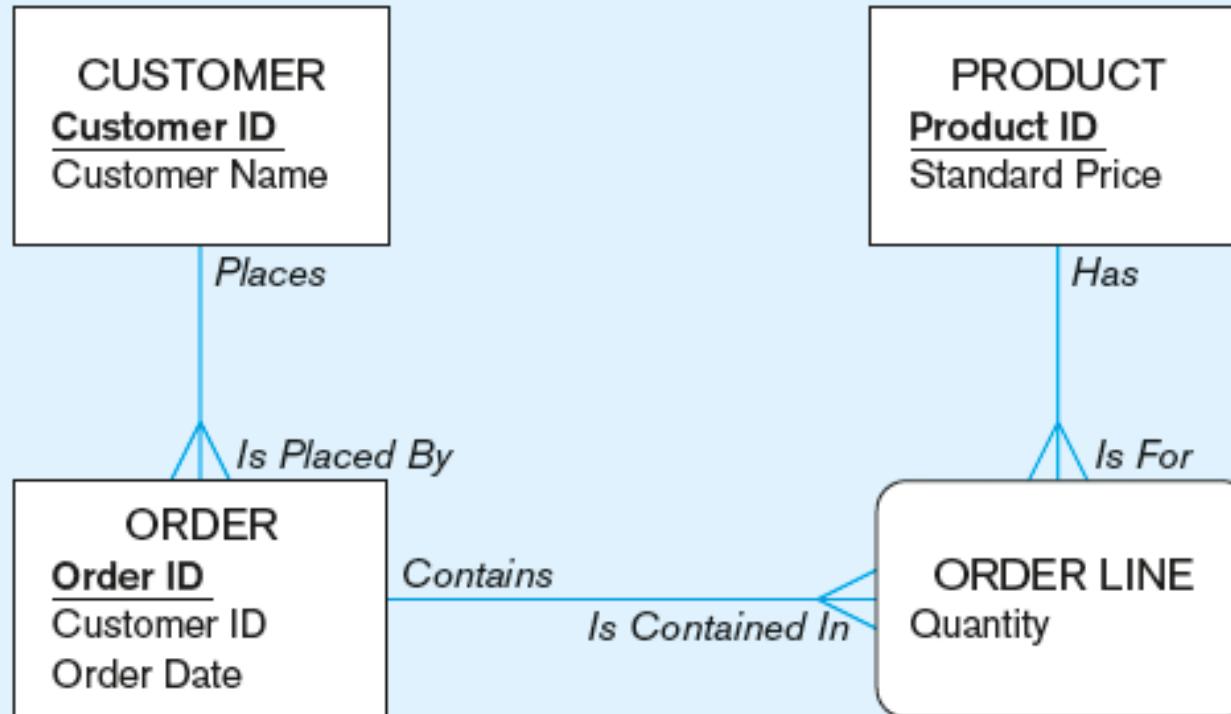
SELF-JOIN

- Definition: A join that matches rows in a table with other rows in the same table.

- It often arises from a unary relationship.

JOIN Examples

The Following Slides will Use This Data Model



Pine Valley Furniture Company Customer_T and Order_T tables with Pointers from Customers to Their Orders

The diagram illustrates the relationship between the Order_T and Customer_T tables. The Order_T table on the left contains records for various orders, each associated with a CustomerID. The Customer_T table on the right lists 16 different furniture stores, each assigned a CustomerID. Arrows connect specific OrderID values in the Order_T table to their corresponding CustomerID values in the Customer_T table. For example, OrderID 1 is connected to CustomerID 1, OrderID 8 to CustomerID 8, and so on. OrderID 0 is connected to a row labeled '(New)'.

Order_T				Customer_T						
	OrderID	OrderDate	CustomerID	CustomerID	CustomerName	CustomerAddress	CustomerCity	CustomerState	CustomerPostalCode	
+	1001	10/21/2010	1	1	Contemporary Casuals	1355 S Hines Blvd	Gainesville	FL	32601-2871	
+	1002	10/21/2010	8	2	Value Furniture	15145 S.W. 17th St.	Plano	TX	75094-7743	
+	1003	10/22/2010	15	3	Home Furnishings	1900 Allard Ave.	Albany	NY	12209-1125	
+	1004	10/22/2010	5	4	Eastern Furniture	1925 Beltline Rd.	Carteret	NJ	07008-3188	
+	1005	10/24/2010	3	5	Impressions	5585 Westcott Ct.	Sacramento	CA	94206-4056	
+	1006	10/24/2010	2	6	Furniture Gallery	325 Flatiron Dr.	Boulder	CO	80514-4432	
+	1007	10/27/2010	11	7	Period Furniture	394 Rainbow Dr.	Seattle	WA	97954-5589	
+	1008	10/30/2010	12	8	California Classics	816 Peach Rd.	Santa Clara	CA	96915-7754	
+	1009	11/5/2010	4	9	M and H Casual Furniture	3709 First Street	Clearwater	FL	34620-2314	
+	1010	11/5/2010	1	10	Seminole Interiors	2400 Rocky Point Dr.	Seminole	FL	34646-4423	
*	0		0	11	American Euro Lifestyles	2424 Missouri Ave N.	Prospect Park	NJ	07508-5621	
*	0		0	12	Battle Creek Furniture	345 Capitol Ave. SW	Battle Creek	MI	49015-3401	
*	0		0	13	Heritage Furnishings	66789 College Ave.	Carlisle	PA	17013-8834	
*	0		0	14	Kaneohe Homes	112 Kiowai St.	Kaneohe	HI	96744-2537	
*	0		0	15	Mountain Scenes	4132 Main Street	Ogden	UT	84403-4432	
*	0		0	(New)						

Record: 1 of 10 | 16 of 16 | No Filter | Search |

EQUI-JOIN

■ Using **WHERE**:

For each customer who placed an order, what is the customer's ID, name and order ID?

```
SELECT Customer_T.CustomerID, CustomerName, OrderID  
FROM Customer_T, Order_T  
WHERE Customer_T.CustomerID = Order_T.CustomerID;
```

■ Using **INNER JOIN ... ON**:

```
SELECT Customer_T.CustomerID, CustomerName, OrderID  
FROM Customer_T INNER JOIN Order_T ON  
Customer_T.CustomerID = Order_T.CustomerID;
```

(Note: In Oracle the keyword **INNER** is optional.)

Outer Join Example

List the customer ID, name, and order number for all customers. Include customer information even for those who have not placed an order.

```
SELECT Customer_T.CustomerID, CustomerName, OrderID  
      FROM Customer_T LEFT OUTER JOIN Order_T ON  
            Customer_T.CustomerID = Order_T.CustomerID;
```

(The LEFT OUTER JOIN clause causes customer data to appear in the result table even if there is no corresponding order data)

The Result: LEFT JOIN

Unlike INNER join, this will include customer rows with no matching order rows

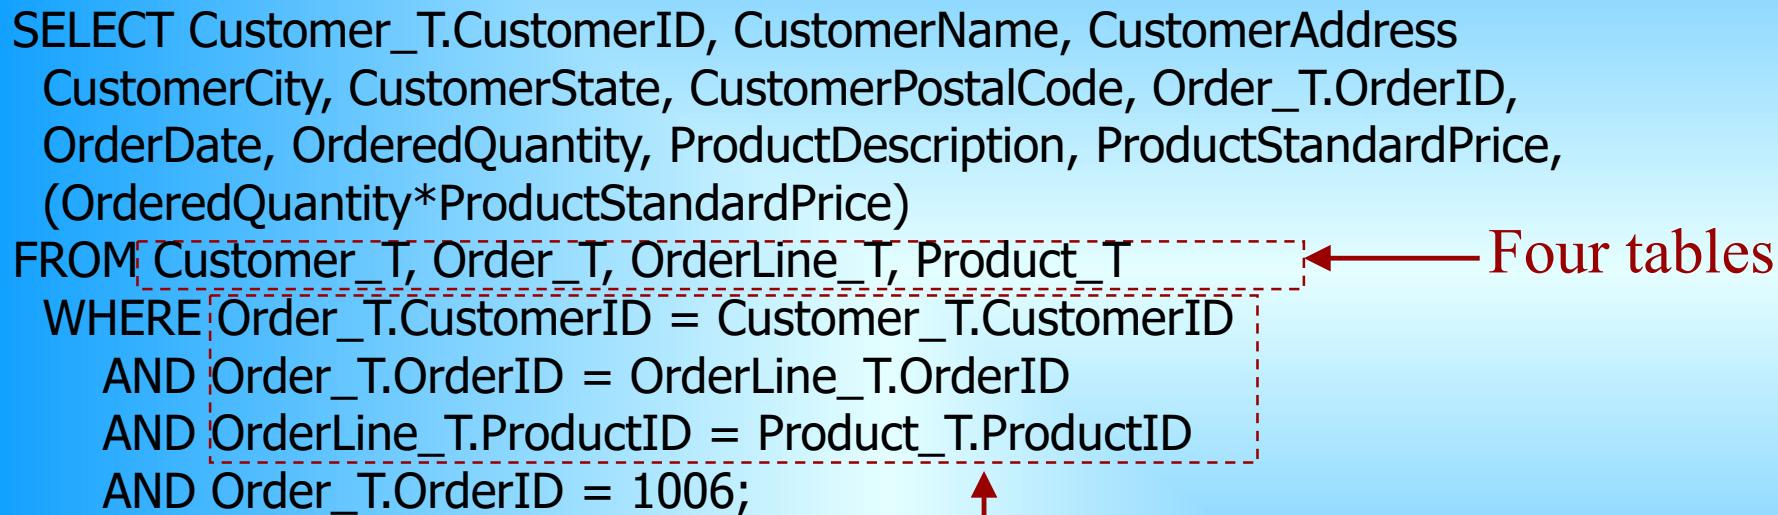
CUSTOMERID	CUSTOMERNAME	ORDERID
1	Contemporary Casuals	1001
8	California Classics	1002
15	Mountain Scenes	1003
5	Impressions	1004
3	Home Furnishings	1005
2	value Furniture	1006
11	American Euro Lifestyles	1007
12	Battle Creek Furniture	1008
4	Eastern Furniture	1009
1	Contemporary Casuals	1010
10	Seminole Interiors	
7	Period Furniture	
14	Kaneohe Homes	
9	M and H Casual Furniture	
13	Heritage Furnishings	
6	Furniture Gallery	

16 rows selected

Multiple Table Join Example

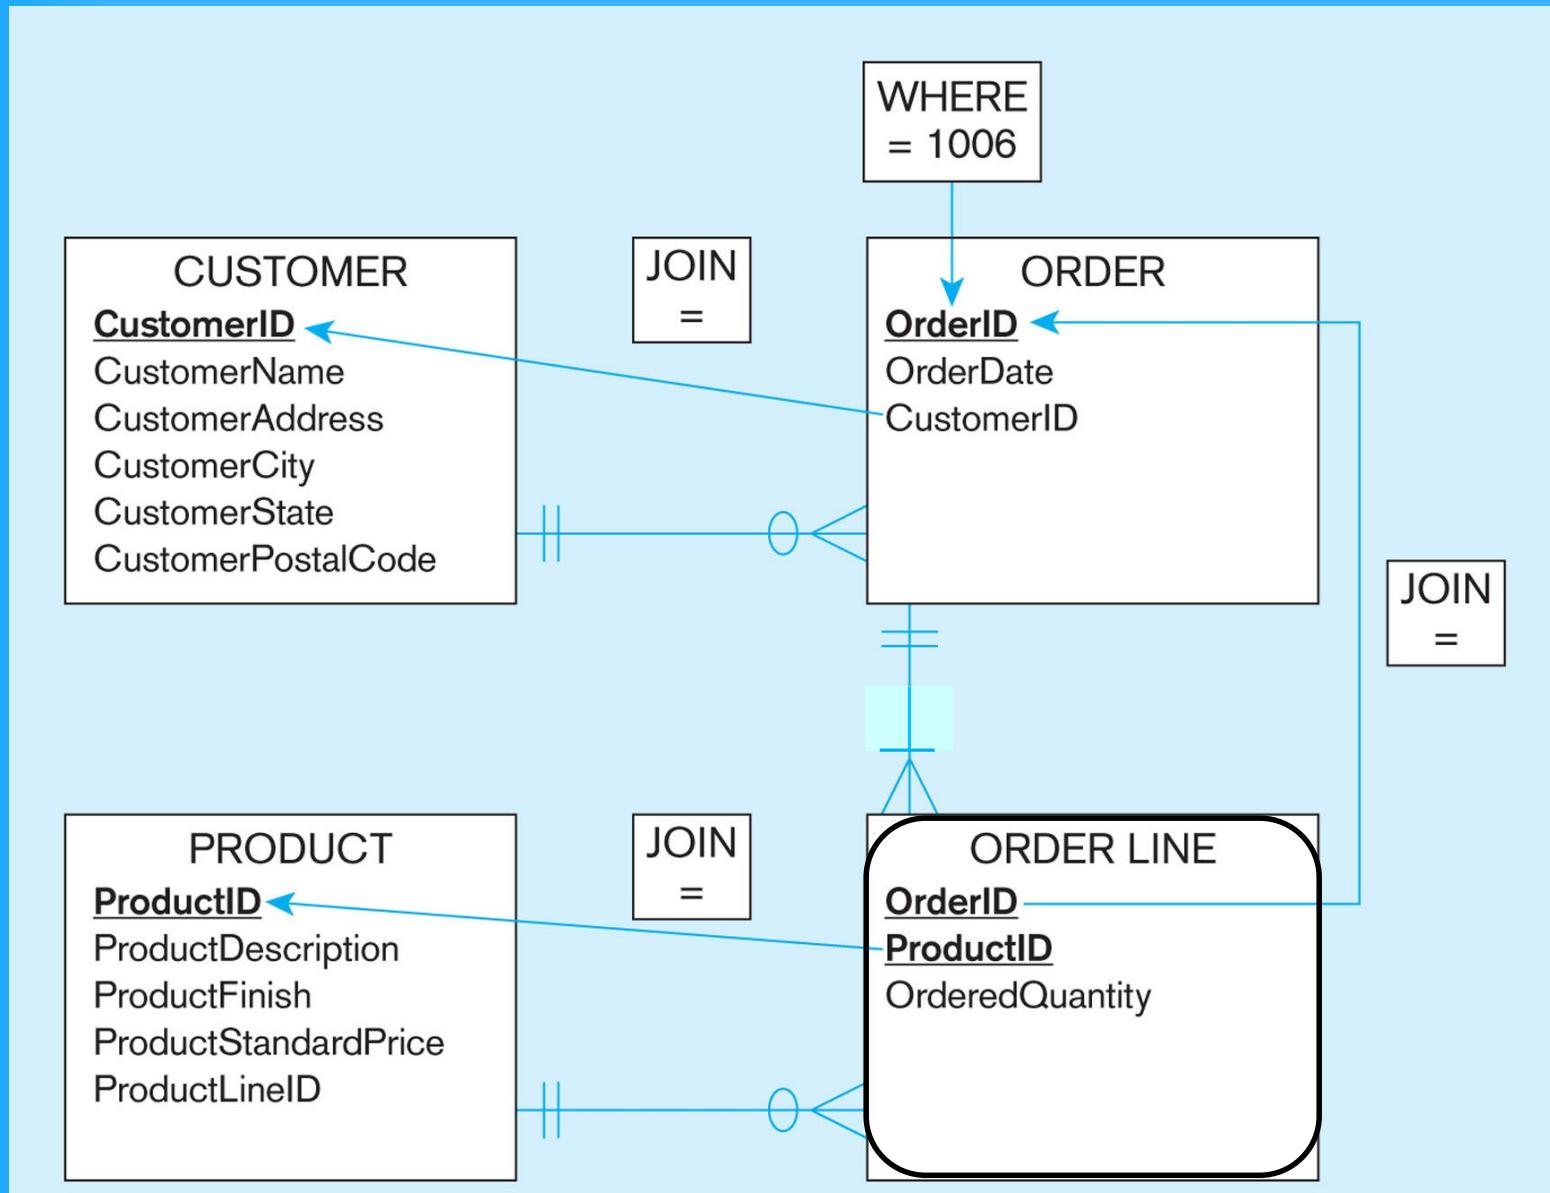
Assemble all information necessary to create an invoice for order number 1006

```
SELECT Customer_T.CustomerID, CustomerName, CustomerAddress  
CustomerCity, CustomerState, CustomerPostalCode, Order_T.OrderID,  
OrderDate, OrderedQuantity, ProductDescription, ProductStandardPrice,  
(OrderedQuantity*ProductStandardPrice)  
FROM Customer_T, Order_T, OrderLine_T, Product_T  
WHERE Order_T.CustomerID = Customer_T.CustomerID  
AND Order_T.OrderID = OrderLine_T.OrderID  
AND OrderLine_T.ProductID = Product_T.ProductID  
AND Order_T.OrderID = 1006;
```



Each pair of tables requires an equality-check condition in the WHERE clause, matching primary keys against foreign keys

Multiple Table Join Example



Multiple Table Join Example

CUSTOMER	
CustomerID
.....
Cx
.....

=

ORDER		
CustomerID	OrderID
Cx	1006
.....
.....

PRODUCT	
ProductID
Px
.....
Py
.....

=

=

ORDER LINE		
ProductID	OrderID
.....
Py	1006
Px	1006
.....

=

=

=

Results from a Four-table Join (Edited for Readability)

From CUSTOMER_T table					
CUSTOMERID	CUSTOMERNAME	CUSTOMERADDRESS	CUSTOMER CITY	CUSTOMER STATE	CUSTOMER POSTALCODE
2	Value Furniture	15145 S. W. 17th St.	Plano	TX	75094 7743
2	Value Furniture	15145 S. W. 17th St.	Plano	TX	75094 7743
2	Value Furniture	15145 S. W. 17th St.	Plano	TX	75094 7743

ORDERID	ORDERDATE	ORDERED QUANTITY	PRODUCTNAME	PRODUCT STANDARDPRICE	(QUANTITY* STANDARDPRICE)
1006	24-OCT -10	1	Entertainment Center	650	650
1006	24-OCT -10	2	Writer's Desk	325	650
1006	24-OCT -10	2	Dining Table	800	1600

From ORDER_T table

From ORDERLINE_T table

From PRODUCT_T table

Self-Join Example

Query: What are the employee ID and name of each employee and the name of his or her supervisor (label the supervisor's name Manager)?

```
SELECT E.EmployeeID, E.EmployeeName, M.EmployeeName AS Manager  
FROM Employee_T E, Employee_T M  
WHERE E.EmployeeSupervisor = M.EmployeeID;
```

Result:

EMPLOYEEID	EMPLOYEENAME	MANAGER
123-44-347	Jim Jason	Robert Lewis

The same table is used on both sides of the join; distinguished using table aliases

Self-joins are usually used on tables with unary relationships

Cartesian Join

(List all possible combinations of the rows of the two tables.
For most times, you want to avoid this.)

- **Missing WHERE clause:**

```
SELECT Customer_T.CustomerID, Order_T.CustomerID,  
CustomerName, OrderID  
FROM Customer_T, Order_T;
```

- **Specifying CROSS JOIN:**

```
SELECT Customer_T.CustomerID, Order_T.CustomerID,  
CustomerName, OrderID  
FROM Customer_T CROSS JOIN Order_T;
```

Processing Multiple Tables Using Subqueries

- **Subquery:** a query (inner query) that is placed inside another query (outer query). Up to **16** levels of nesting are typically supported.
- Options:
 - In a condition of the WHERE clause
 - As a "table" of the FROM clause
 - Within the HAVING clause
- Two kinds of subqueries:
 - **Correlated:** executed once for each row returned by the outer query
 - **Noncorrelated:** executed once for the entire outer query

Subquery in the WHERE Clause

Show all customers who have placed an order

```
SELECT CustomerName  
      FROM Customer_T  
     WHERE CustomerID IN  
           (SELECT DISTINCT CustomerID  
            FROM Order_T);
```



Subquery is embedded in parentheses.
In this case it returns a list that will be
used in the WHERE clause of the outer
query

The IN operator will
test to see if the
CustomerID value of a
row is included in the
list returned from the
subquery

Result:

CUSTOMER_NAME
Contemporary Casuals
Value Furniture
Home Furnishings
Eastern Furniture
Impressions
California Classics
American Euro Lifestyles
Battle Creek Furniture
Mountain Scenes

9 rows selected.

Subquery in the FROM Clause

Show all products whose standard price is higher than the average price

Subquery forms the derived table used in the FROM clause of the outer query

One column of the subquery is an aggregate function that has an alias name. That alias can then be referred to in the outer query

```
SELECT ProductDescription, ProductStandardPrice, AvgPrice  
FROM  
  (SELECT AVG(ProductStandardPrice) AvgPrice FROM Product_T),  
    Product_T  
WHERE ProductStandardPrice > AvgPrice;
```

The WHERE clause normally cannot include aggregate functions, but because the aggregate is performed in the subquery its result can be used in the outer query's WHERE clause

Subquery Within the HAVING Clause

Find all the types of product finishes and their frequencies in the Product_T table except the most popular finish.

```
SELECT COUNT(*), ProductFinish  
FROM Product_T  
GROUP BY ProductFinish  
Having COUNT(*) <  
    (SELECT MAX(COUNT(*))  
     FROM Product_T  
     GROUP BY ProductFinish);
```

Note: The HAVING clause works together with GROUP BY clause.

Correlated vs. Noncorrelated Subqueries

■ **Noncorrelated subqueries:**

- Do not depend on data from the outer query
- Execute once for the entire outer query

■ **Correlated subqueries:**

- Depends on data from the outer query
- Execute once for each row of the outer query
- Often use the **EXISTS** operator in the WHERE clause of the outer query

Processing a Noncorrelated Subquery

What are the names of customers who have placed orders?

```
SELECT CustomerName  
      FROM Customer_T  
     WHERE CustomerID IN
```

```
(SELECT DISTINCT CustomerID  
      FROM Order_T);
```

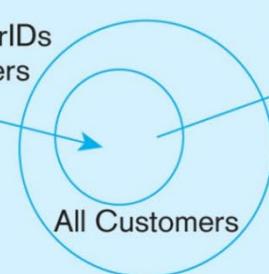
1. The subquery (shown in the box) is processed first and an intermediate results table created:

CUSTOMERID

1
8
15
5
3
2
11
12
4

9 rows selected.

CustomerIDs
from orders



2. The outer query returns the requested customer information for each customer included in the intermediate results table:

CUSTOMERNAME

Contemporary Casuals
Value Furniture
Home Furnishings
Eastern Furniture
Impressions
California Classics
American Euro Lifestyles
Battle Creek Furniture
Mountain Scenes
9 rows selected.

A noncorrelated subquery processes completely before the outer query begins, i.e., the subqueries are processed **from the inside out**.

Correlated Subquery Example

Show all orders that include furniture finished in natural ash

```
SELECT DISTINCT OrderID FROM OrderLine_T
WHERE EXISTS <--  

  (SELECT * FROM Product_T  

   WHERE ProductID = OrderLine_T.ProductID  

    AND Productfinish = 'Natural ash');
```

The EXISTS operator will return a TRUE value if the subquery resulted in a non-empty set, otherwise it returns a FALSE; this keyword is often an indication of a correlated subquery

The subquery is testing for a value that comes from the outer query

Note: A correlated subquery always refers to an attribute from a table referenced in the outer query; a correlated subquery is processed **from the outside in**

Processing a Correlated Subquery

Subquery refers to outer-query data, so executes once for each row of outer query

What are the order IDs for all orders that have included furniture finished in natural ash?

```
SELECT DISTINCT OrderID FROM OrderLine_T
```

WHERE EXISTS

(SELECT *

FROM Product_T

WHERE ProductID = OrderLine_T.ProductID
AND ProductFinish = 'Natural Ash');

	OrderID	ProductID	OrderedQuantity
1	1001	1	1
	1001	2	2
	1001	4	1
	1002	3	5
3	1003	3	3
	1004	6	2
	1004	8	2
	1005	4	4
	1006	4	1
	1006	5	2
	1007	1	3
	1007	2	2
	1008	3	3
	1008	8	3
	1009	4	2
	1009	7	3
	1010	8	10
*	0	0	0

	ProductID	ProductDescription	ProductFinish	ProductStandardPrice	ProductLineID
►	1	End Table	Cherry	\$175.00	10001
►	2	Coffee Table	Natural Ash	\$200.00	20001
►	4	Computer Desk	Natural Ash	\$375.00	20001
►	4	Entertainment Center	Natural Maple	\$650.00	30001
►	5	Writer's Desk	Cherry	\$325.00	10001
►	6	8-Drawer Dresser	White Ash	\$750.00	20001
►	7	Dining Table	Natural Ash	\$800.00	20001
►	8	Computer Desk	Walnut	\$250.00	30001
*	(AutoNumber)			\$0.00	

Note: only the orders that involve products with Natural Ash will be included in the final results

1. The first order ID is selected from OrderLine_T: OrderID =1001.
2. The subquery is evaluated to see if any product in that order has a natural ash finish. Product 2 does, and is part of the order. EXISTS is valued as *true* and the order ID is added to the result table.
3. The next order ID is selected from OrderLine_T: OrderID =1002.
4. The subquery is evaluated to see if the product ordered has a natural ash finish. It does. EXISTS is valued as *true* and the order ID is added to the result table.

ORDERID

1001
1002
1003
1007
1008
1009

Union Queries

Combining the outputs (union of multiple queries) together into a single result table

- Queries involved must output the same number of columns
- An output from each query for each column must have compatible data types
- It is safest to use the CAST command to control the data type conversion, e.g.,

```
SELECT CAST (OrderDate AS CHAR) FROM Order_T;
```

Union Query Example

You just
need to
define
once for a
column
name

To combine
two query
results

```
SELECT C1.CustomerID, CustomerName, OrderedQuantity,  
'Largest Quantity' AS Quantity  
FROM Customer_T C1,Order_T O1, OrderLine_T Q1  
WHERE C1.CustomerID = O1.CustomerID  
AND O1.OrderID = Q1.OrderID  
AND OrderedQuantity =  
(SELECT MAX(OrderedQuantity)  
FROM OrderLine_T)
```

First query

UNION

```
SELECT C1.CustomerID, CustomerName, OrderedQuantity,  
'Smallest Quantity'  
FROM Customer_T C1, Order_T O1, OrderLine_T Q1  
WHERE C1.CustomerID = O1.CustomerID  
AND O1.OrderID = Q1.OrderID  
AND OrderedQuantity =  
(SELECT MIN(OrderedQuantity)  
FROM OrderLine_T)
```

Second query

ORDER BY 3;

Combining Queries Using UNION

```
SELECT C1.CustomerID, CustomerName, OrderedQuantity, 'Largest Quantity' AS Quantity
  FROM Customer_T C1,Order_T O1, OrderLine_T Q1
 WHERE C1.CustomerID = O1.CustomerID
       AND O1.OrderID = Q1.OrderID
       AND OrderedQuantity =
            (SELECT MAX(OrderedQuantity)
              FROM OrderLine_T)
```

1. In the above query, the subquery is processed first and an intermediate results table created. It contains the maximum quantity ordered from OrderLine_T and has a value of 10.
2. Next the main query selects customer information for the customer or customers who ordered 10 of any item. Contemporary Casuals has ordered 10 of some unspecified item.

Note: with UNION queries, the quantity and data types of the attributes in the SELECT clauses of both queries must be identical

```
SELECT C1.CustomerID, CustomerName, OrderedQuantity, 'Smallest Quantity'
  FROM Customer_T C1, Order_T O1, OrderLine_T Q1
 WHERE C1.CustomerID = O1.CustomerID
       AND O1.OrderID = Q1.OrderID
       AND OrderedQuantity =
            (SELECT MIN(OrderedQuantity)
              FROM OrderLine_T)

ORDER BY 3;
```

1. In the second main query, the same process is followed but the result returned is for the minimum order quantity.
2. The results of the two queries are joined together using the UNION command.
3. The results are then ordered according to the value in OrderedQuantity. The default is ascending value, so the orders with the smallest quantity, 1, are listed first.

The Result of the Example

CUSTOMERID	CUSTOMERNAME	ORDEREDQUANTUTY	QUANTUTY
1	Contemporary Casuals	1	Smallest Quantity
2	Value Furniture	1	Smallest Quantity
1	Contemporary Casuals	10	Largest Quantity

Some More Complicated SQL Queries

Question: For each salesperson, list his or her biggest selling product.

```
CREATE VIEW TSales_V AS
SELECT SalespersonName, ProductDescription,
       SUM(OrderedQuantity) Totorders
FROM Salesperson_T, OrderLine_T, Product_T, Order_T
  WHERE Salesperson_T.SalespersonID=Order_T.SalespersonID
    AND Order_T.OrderID=OrderLine_T.OrderID
    AND OrderLine_T.ProductID=Product_T.ProductID
 GROUP By SalespersonName, ProductDescription;
```

```
SELECT SalespersonName, ProductDescription
  FROM Tsales_V A
 WHERE Totorders=(SELECT MAX(Totorders)
   FROM Tsales_V B
 WHERE B.SalespersonName=
      A.SalespersonName);
```

Guidelines for Query Design

- Understand the desired results, especially the attributes desired in results.
- Be familiar with the data model (entities and relationships).
- Review ERD.
- Identify the entities that contain desired attributes.
- Construct a WHERE equality for each link.
- Fine tune with GROUP BY and HAVING clauses if needed.
- Consider the effect on unusual data.

Guidelines for Query Design (cont.)

- Understand how indexes are used in query processing.
- Keep optimizer statistics up-to-date (some DBMSes do it automatically but costly).
- Use compatible data types for fields and literals (in queries).
- Write simple queries (can use UNION to combine).
- Break complex queries into multiple simple parts.
- Try not to nest one query inside another query.
- Try not to combine a query with itself (self-joins) but use a temporary copy of the table.

Guidelines for Query Design (cont.)

- Create temporary tables for groups of queries.
`(CREATE GLOBAL TEMPORARY TABLE tempTable AS ...;)`
- Combine update operations (DBMS may do it in parallel).
- Retrieve only the data you need (do not use `SELECT *` when not retrieving all the fields).
- Try to avoid sorting without an index on the sort key field (often faster to do it outside a DBMS).
- Learn by tracking your DBMS process data (with the `EXPLAIN PLAN FOR` command).
- Consider the total time needed for ad hoc queries (just write logically correct queries).

Create a Temporary Table

Syntax:

```
CREATE GLOBAL TEMPORARY TABLE tempTableName  
    ON COMMIT PRESERVE ROWS  
    AS SELECT * FROM realTable;
```

Example:

```
CREATE GLOBAL TEMPORARY TABLE tempCustomer_t  
    ON COMMIT PRESERVE ROWS  
    AS SELECT * FROM customer_t;
```

Output:

Global temporary TABLE created.

Ensuring Transaction Integrity

- **Transaction:** A single unit of work that must be completely processed or not processed at all
 - It may involve multiple updates.
 - If any update fails, then all other updates must be cancelled.
- SQL commands for transactions
 - BEGIN TRANSACTION/END TRANSACTION
 - Marks boundaries of a transaction
 - COMMIT
 - Makes all updates permanent
 - ROLLBACK
 - Cancels updates since the last COMMIT

An SQL Transaction Sequence (in Pseudocode)

Different DBMS
may have different syntaxes.

BEGIN transaction

INSERT OrderID, Orderdate, CustomerID into Order_T;

INSERT OrderID, ProductID, OrderedQuantity into OrderLine_T;
INSERT OrderID, ProductID, OrderedQuantity into OrderLine_T;
INSERT OrderID, ProductID, OrderedQuantity into OrderLine_T;

END transaction

Either all the rows are entered together or none of them should be entered; it is all or nothing at all.

Valid information inserted.
COMMIT work.

All changes to data
are made permanent.

Invalid ProductID entered.

Transaction will be ABORTED.
ROLLBACK all changes made to Order_T.

All changes made to Order_T
and OrderLine_T are removed.
Database state is just as it was
before the transaction began.

An Oracle Example

Create and populate a table:

```
CREATE TABLE Mytable_T  
  (id      NUMBER    NOT NULL,  
   description  VARCHAR2(50) NOT NULL);
```

```
INSERT INTO Mytable_T (id, description) VALUES (1, 'Description for 1');  
INSERT INTO Mytable_T (id, description) VALUES (2, 'Description for 2');
```

```
BEGIN  
  FOR i IN 3 .. 10 LOOP  
    INSERT INTO Mytable_T (id, description)  
    VALUES (i, 'Description for ' || i);  
  END LOOP;  
END;  
/
```

Check table contents:

```
SELECT * FROM Mytable_T;
```

ID	DESCRIPTION
1	Description for 1
2	Description for 2
3	Description for 3
4	Description for 4
5	Description for 5
6	Description for 6
7	Description for 7
8	Description for 8
9	Description for 9
10	Description for 10

An Oracle Example (cont.)

Create and populate a table with ROLLBACK:

```
CREATE TABLE Mytable_T  
  (id      NUMBER      NOT NULL,  
   description  VARCHAR2(50) NOT NULL);  
  
INSERT INTO Mytable_T (id, description) VALUES (1, 'Description for 1');  
INSERT INTO Mytable_T (id, description) VALUES (2, 'Description for 2');  
  
BEGIN  
  FOR i IN 3 .. 10 LOOP  
    INSERT INTO Mytable_T (id, description)  
    VALUES (i, 'Description for ' || i);  
  END LOOP;  
ROLLBACK;  
END;  
/
```

The ROLLBACK statement removes all the insertions

Check table contents:

```
SELECT * FROM Mytable_T;
```

ID	DESCRIPTION

An Oracle Example (cont.)

Using **SAVEPOINT** and **ROLLBACK**:

```
CREATE TABLE Mytable_T
  (id      NUMBER      NOT NULL,
   description  VARCHAR2(50) NOT NULL);

INSERT INTO Mytable_T (id, description) VALUES (1, 'Description for 1');
INSERT INTO Mytable_T (id, description) VALUES (2, 'Description for 2');

SAVEPOINT a;
BEGIN
  FOR i IN 3 .. 10 LOOP
    INSERT INTO Mytable_T (id, description)
      VALUES (i, 'Description for ' || i);
  END LOOP;
ROLLBACK to a;
END;
/
```

The **ROLLBACK** statement removes all the insertions up to **SAVEPOINT a**

Check table contents:

```
SELECT * FROM Mytable_T;
```

ID	DESCRIPTION
1	Description for 1
2	Description for 2

An Oracle Example (cont.)

Using two SAVEPOINTS and ROLLBACK:

```
CREATE TABLE Mytable_T
  (id      NUMBER      NOT NULL,
   description  VARCHAR2(50) NOT NULL);

INSERT INTO Mytable_T (id, description) VALUES (1, 'Description for 1');
SAVEPOINT a;
INSERT INTO Mytable_T (id, description) VALUES (2, 'Description for 2');
SAVEPOINT b;
BEGIN
  FOR i IN 3 .. 10 LOOP
    INSERT INTO Mytable_T (id, description)
    VALUES (i, 'Description for ' || i);
  END LOOP;
ROLLBACK to a;
END;
/
```

You can use many
SAVEPOINTS; the
ROLLBACK statement
can roll back to a
specific SAVEPOINT

Check table contents:

```
SELECT * FROM Mytable_T;
```

ID	DESCRIPTION
1	Description for 1

An Oracle Example (cont.)

COMMIT and then ROLLBACK:

```
CREATE TABLE Mytable_T  
  (id      NUMBER      NOT NULL,  
   description  VARCHAR2(50) NOT NULL);
```

```
INSERT INTO Mytable_T (id, description) VALUES (1, 'Description for 1');  
INSERT INTO Mytable_T (id, description) VALUES (2, 'Description for 2');  
SAVEPOINT a;
```

```
BEGIN  
  FOR i IN 3 .. 10 LOOP  
    INSERT INTO Mytable_T (id, description)  
    VALUES (i, 'Description for ' || i);  
  END LOOP;  
COMMIT;  
ROLLBACK ;  
END;  
/
```

After issuing the COMMIT command, the insertions became final and the ROLLBACK statement cannot remove the insertions

ID	DESCRIPTION
1	Description for 1
2	Description for 2
3	Description for 3
4	Description for 4
5	Description for 5
6	Description for 6
7	Description for 7
8	Description for 8
9	Description for 9
10	Description for 10

Check table contents:

```
SELECT * FROM Mytable_T;
```

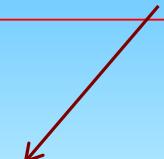
Data Dictionary Facilities

- **Data Dictionary:** System tables that store metadata (DB definition info), valuable info for users and DBAs.
- Users usually can view some of these tables by using SELECT statements.
- Users are restricted from updating them.
- Some examples in Oracle 11g (which has a total of 522 views for DBAs). You must be logged in as a DBA.
 - DBA_TABLES – descriptions of tables.
 - DBA_CONSTRAINTS – descriptions of constraints.
 - DBA_USERS – information about the users of the system.
- Examples in Microsoft SQL Server 2008
 - sys.columns – table and column definitions.
 - sys.indexes – table index information.
 - sys.foreign_key_columns – details about columns in foreign key constraints.

Some Oracle Examples

Find out the owner of Product_T table:

```
SELECT OWNER, TABLE_NAME  
      FROM DBA_TABLES  
     WHERE TABLE_NAME = 'PRODUCT_T';
```



	OWNER	TABLE_NAME
1	SYSTEM	PRODUCT_T
2	STUDENT1	PRODUCT_T
3	XY	PRODUCT_T
4	YJN	PRODUCT_T

Some Oracle Examples (cont.)

Find out all the information about all the tables in the database

```
SELECT *
FROM ALL_TABLES;
```

OWNER	TABLE_NAME	TABLESPACE_NAME	CLUSTER_NAME	IOT_NAME	STATUS	PCT_FREE	PCT_USED	INI_TRANS	MAX_TRANS
1 SYS	ICOL\$	SYSTEM	C_OBJ#	(null)	VALID	0	0	0	0
2 SYS	CON\$	SYSTEM	(null)	(null)	VALID	10	40	1	255
3 SYS	UNDO\$	SYSTEM	(null)	(null)	VALID	10	40	1	255
4 SYS	PROXY_ROLE_DATA\$	SYSTEM	(null)	(null)	VALID	10	40	1	255
5 SYS	FILE\$	SYSTEM	(null)	(null)	VALID	10	40	1	255
6 SYS	UET\$	SYSTEM	C_FILE#_BLOCK#	(null)	VALID	0	0	0	0
7 SYS	IND\$	SYSTEM	C_OBJ#	(null)	VALID	0	0	0	0
8 SYS	SEG\$	SYSTEM	C_FILE#_BLOCK#	(null)	VALID	0	0	0	0
9 SYS	COL\$	SYSTEM	C_OBJ#	(null)	VALID	0	0	0	0
10 SYS	CLU\$	SYSTEM	C_OBJ#	(null)	VALID	0	0	0	0
11 SYS	PROXY_DATA\$	SYSTEM	(null)	(null)	VALID	10	40	1	255
12 SYS	TS\$	SYSTEM	C_TS#	(null)	VALID	0	0	0	0
13 SYS	BOOTSTRAP\$	SYSTEM	(null)	(null)	VALID	10	40	1	255
14 SYS	FET\$	SYSTEM	C_TS#	(null)	VALID	0	0	0	0
15 SYS	CCOL\$	SYSTEM	C_COBJ#	(null)	VALID	0	0	0	0
16 SYS	USER\$	SYSTEM	C_USER#	(null)	VALID	0	0	0	0
17 SYS	OBJ\$	SYSTEM	(null)	(null)	VALID	10	40	1	255
18 SYS	TAB\$	SYSTEM	C_OBJ#	(null)	VALID	0	0	0	0
19 SYS	CDEF\$	SYSTEM	C_COBJ#	(null)	VALID	0	0	0	0
20 SYS	OBJERROR\$	SYSTEM	(null)	(null)	VALID	10	40	1	255
21 SYS	OBJAUTH\$	SYSTEM	(null)	(null)	VALID	10	40	1	255
..									

Some Oracle Examples (cont.)

Find out all the users of the database:

```
SELECT *
  FROM DBA_USERS;
```

	USERNAME	USER_ID	PASSWORD	ACCOUNT_STATUS	LOCK_DATE	EXPIRY_DATE	DEFAULT_TABLESPACE	TEMPORARY_TABLESPACE
1	SYSTEM	5	(null)	OPEN	(null)	15-JUN-12	SYSTEM	TEMP
2	SYS	0	(null)	OPEN	(null)	15-JUN-12	SYSTEM	TEMP
3	YJN	55	(null)	OPEN	(null)	21-AUG-12	SYSTEM	TEMP
4	MARVIN	49	(null)	OPEN	(null)	16-JUN-12	SYSTEM	TEMP
5	STUDENT1	54	(null)	OPEN	(null)	15-JUL-12	SYSTEM	TEMP
6	ANONYMOUS	35	(null)	OPEN	(null)	23-FEB-12	SYSAUX	TEMP
7	HR	43	(null)	OPEN	(null)	17-JUN-12	USERS	TEMP
8	XY	48	(null)	OPEN	(null)	15-JUN-12	USERS	TEMP
9	XIAOLONG	50	(null)	LOCKED (TIMED)	10-JAN-12	16-JUN-12	USERS	TEMP
10	APEX_PUBLIC_USER	45	(null)	LOCKED	27-AUG-11	23-FEB-12	SYSTEM	TEMP
11	APEX_040000	47	(null)	LOCKED	27-AUG-11	23-FEB-12	SYSAUX	TEMP
12	FLOWS_FILES	44	(null)	LOCKED	27-AUG-11	23-FEB-12	SYSAUX	TEMP
13	XS\$NULL	2147483638	(null)	EXPIRED & LOCKED	27-AUG-11	27-AUG-11	SYSTEM	TEMP
14	OUTLN	9	(null)	EXPIRED & LOCKED	18-DEC-11	18-DEC-11	SYSTEM	TEMP
15	XDB	34	(null)	EXPIRED & LOCKED	27-AUG-11	27-AUG-11	SYSAUX	TEMP
16	CTXSYS	32	(null)	EXPIRED & LOCKED	18-DEC-11	18-DEC-11	SYSAUX	TEMP
17	MDSYS	42	(null)	EXPIRED & LOCKED	27-AUG-11	18-DEC-11	SYSAUX	TEMP

SQL:1999 and SQL:200N Enhancements/Extensions

- User-defined data types (UDT)
 - Users can define subclasses of standard types or an object type, which may contain functions/methods
- Analytical functions (for OLAP)
 - CEILING, FLOOR, SQRT, RANK, DENSE_RANK, ROLLUP, CUBE, SAMPLE,
 - WINDOW—improved numerical analysis capabilities
- New Data Types
 - BIGINT, MULTISET (collection), XML
- CREATE TABLE LIKE—create a new table similar to an existing one
- MERGE

SQL:1999 and SQL:200N Enhancements/Extensions (cont.)

- Persistent Stored Modules (SQL/PSM)
 - Capability to create and drop code modules
 - New statements:
 - CASE, IF, LOOP, FOR, WHILE, etc.
 - Makes SQL into a **procedural language**
- Oracle has proprietary version called **PL/SQL**, and Microsoft SQL Server has Transact/SQL

Oracle PL/SQL

PL/SQL stands for Procedural Language extension of SQL.

PL/SQL is a combination of SQL along with the procedural features of programming languages. It was developed by Oracle Corporation in the early 1990's to enhance the capabilities of SQL.

A PL/SQL Block consists of three sections:

The Declaration section (optional).

The Execution section (mandatory).

The Exception (or Error) Handling section (optional).

DECLARE
Variable declaration
BEGIN
Program Execution
EXCEPTION
Exception handling
END;

An Oracle Example: Using RANK()

Find the employee with the highest salary:

```
SELECT *
FROM (
    SELECT employee_id, last_name, salary,
RANK() OVER (ORDER BY salary DESC) EMPRANK
    FROM employees)
WHERE EMPRANK = 1;
```

EMPLOYEE_ID	LAST_NAME	SALARY	EMPRANK
1	Yang	2000000	1
2	Shen	2000000	1

An Oracle Example: Using PL/SQL

Calculate bonus for an employee:

```
DECLARE
    bonus NUMBER(8,2);
    emp_id NUMBER(6) := 1;
BEGIN
    SELECT salary * 0.10 INTO bonus
        FROM employees
        WHERE employee_id = emp_id;
    DBMS_OUTPUT.PUT_LINE (bonus);
END;
/
```

PL/SQL procedure successfully completed.
200000

An Oracle Example: Using WHILE and LOOP Statements in PL/SQL

```
DECLARE
```

```
    done BOOLEAN; -- Initialize to NULL by default  
    counter NUMBER := 0;
```

```
BEGIN
```

```
    done := FALSE; -- Assign literal value  
    WHILE done != TRUE -- Compare to literal value  
        LOOP
```

```
        counter := counter + 1;  
        done := (counter > 500); -- Assign value of BOOLEAN expression  
    END LOOP;
```

```
    DBMS_OUTPUT.PUT_LINE ('Counter = '|| counter);
```

```
END;
```

```
/
```

PL/SQL procedure successfully completed.

Counter = 501

An Oracle Example: Using CASE Statement in PL/SQL

```
DECLARE
```

```
    grade CHAR(1) := 'B';
    appraisal VARCHAR2(20);
```

```
BEGIN
```

```
    appraisal :=
        CASE grade
            WHEN 'A' THEN 'Excellent'
            WHEN 'B' THEN 'Very Good'
            WHEN 'C' THEN 'Good'
            WHEN 'D' THEN 'Fair'
            WHEN 'F' THEN 'Poor'
            ELSE 'No such grade'
        END;
```

```
DBMS_OUTPUT.PUT_LINE
```

```
('Grade ' || grade || ' is ' || appraisal);
```

```
END;
```

```
/
```

PL/SQL procedure successfully completed.

Grade B is Very Good

Routines and Triggers

- **Routines:**

Stored program modules (or subprograms) that execute on demand (or must be called explicitly)

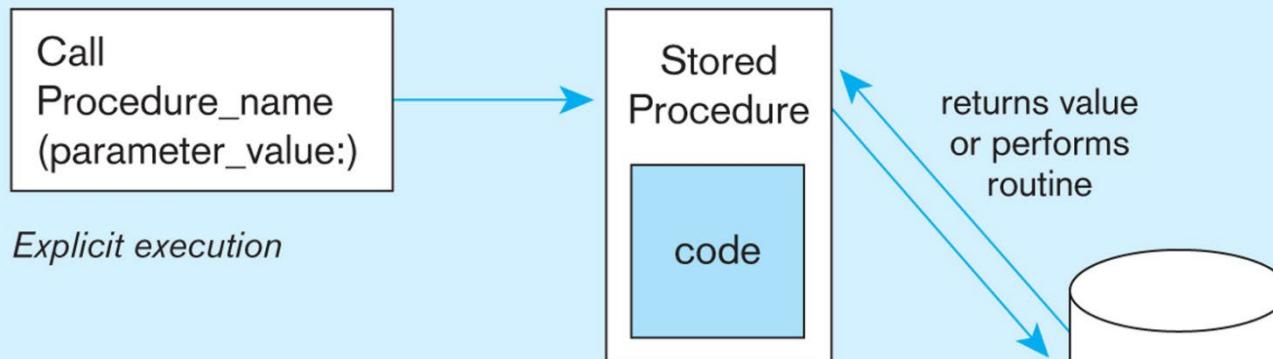
- **Function:** a routine that returns a value

- **Procedure:** a routine that performs a specific action but does not return a value

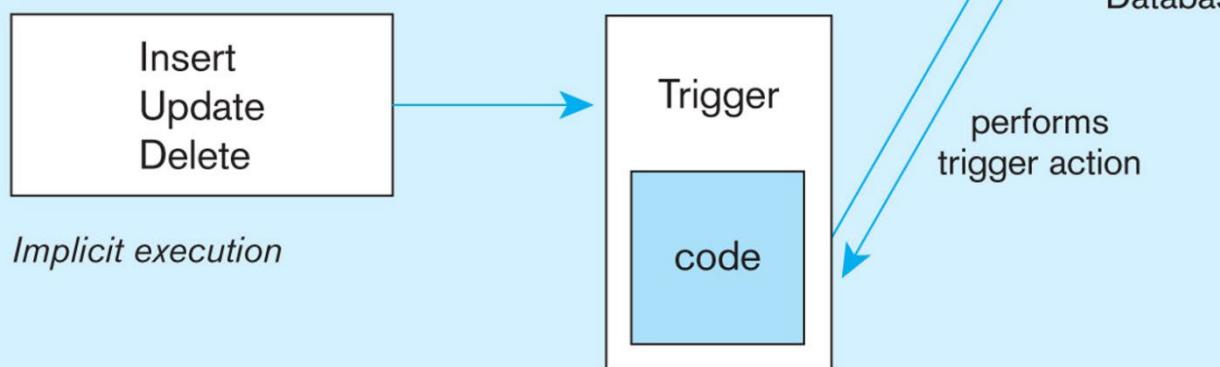
- **Trigger:** a named set of SQL statements that executes automatically in response to a database event (e.g., INSERT, UPDATE, or DELETE)

Triggers Contrasted with Stored Routines

ROUTINE: Procedures and functions are called explicitly



TRIGGER: Triggers are event-driven



Source: adapted from Mullins, 1995.

Syntax for Creating a Procedure or Function, SQL:200n

```
{CREATE PROCEDURE | CREATE FUNCTION} routine_name  
([parameter [{,parameter} . . .]])  
[RETURNS data_type result_cast] /* for functions only */  
[LANGUAGE {ADA | C | COBOL | FORTRAN | MUMPS | PASCAL | PLI | SQL}]  
[PARAMETER STYLE {SQL | GENERAL}]  
[SPECIFIC specific_name]  
[DETERMINISTIC | NOT DETERMINISTIC]  
[NO SQL | CONTAINS SQL | READS SQL DATA | MODIFIES SQL DATA]  
[RETURNS NULL ON NULL INPUT | CALLED ON NULL INPUT]  
[DYNAMIC RESULT SETS unsigned_integer] /* for procedures only */  
[STATIC DISPATCH] /* for functions only */  
[NEW SAVEPOINT LEVEL | OLD SAVEPOINT LEVEL]  
routine_body
```

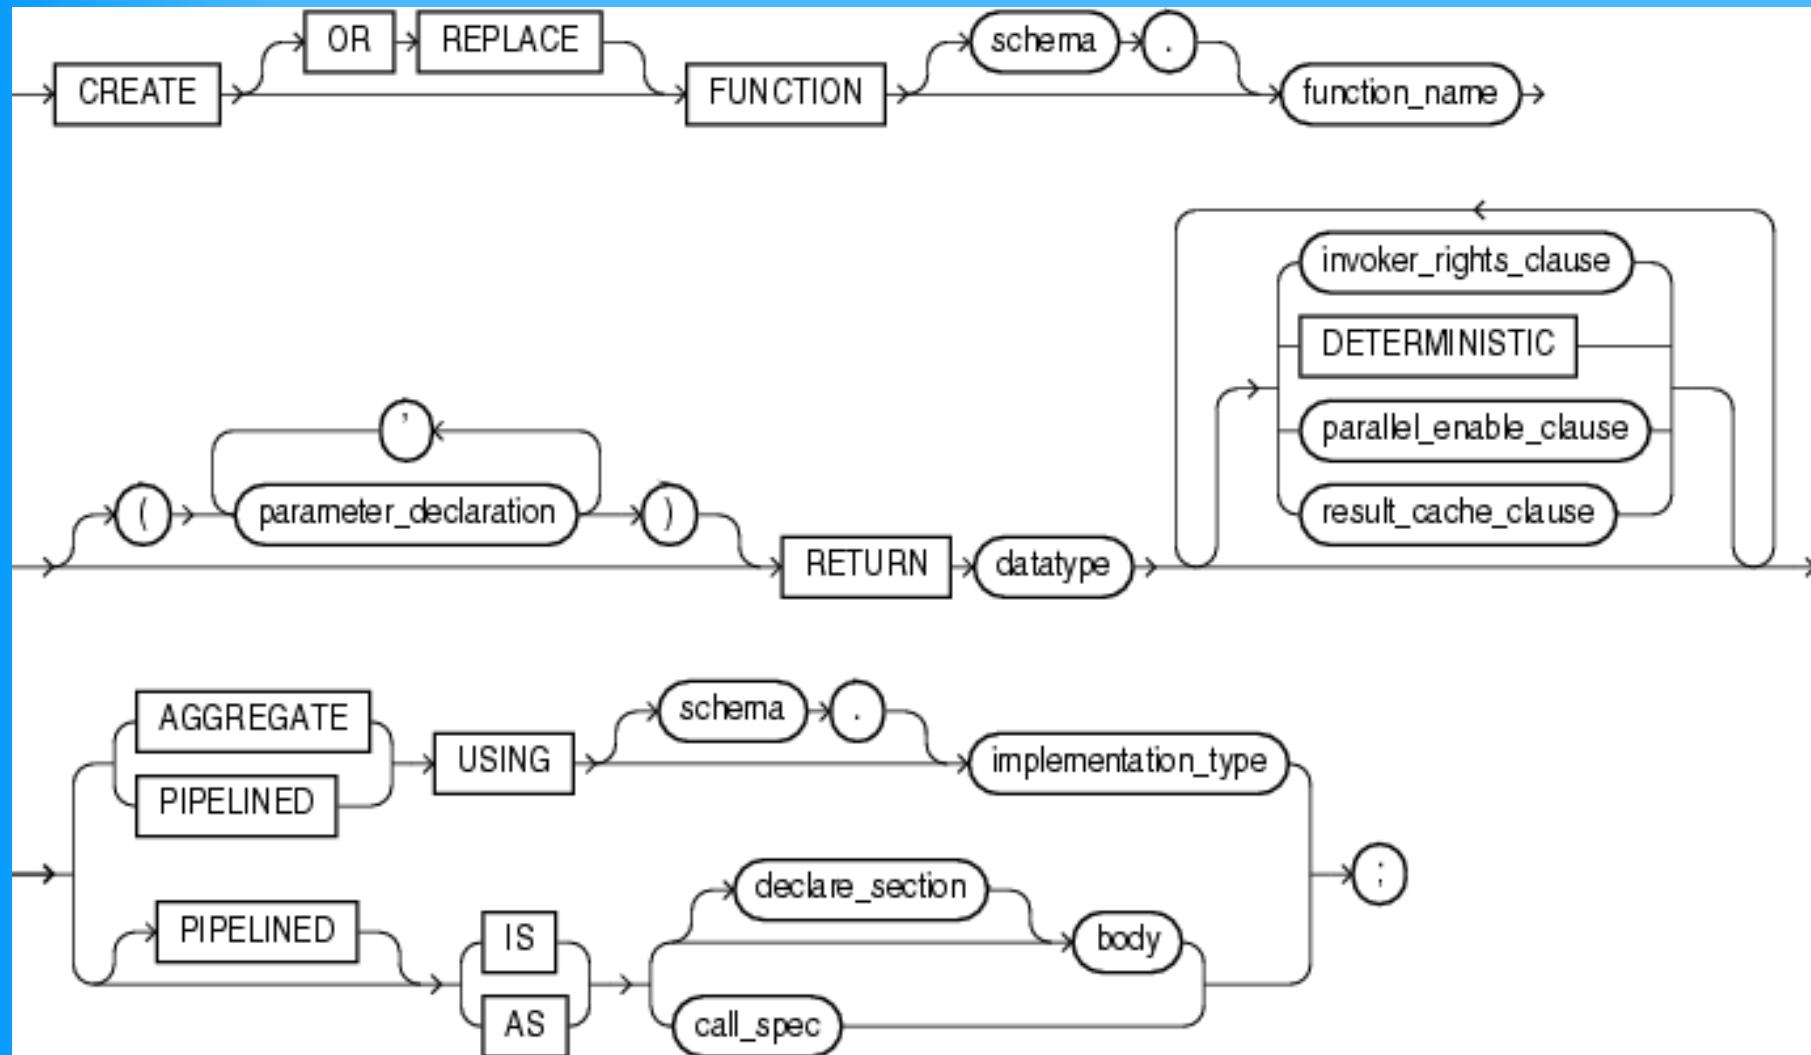
TABLE 7-2 Comparison of Vendor Syntax Differences in Stored Procedures

The vendors' syntaxes differ in stored procedures more than in ordinary SQL. For an illustration, here is a chart that shows what CREATE PROCEDURE looks like in three dialects. We use one line for each significant part, so you can compare dialects by reading across the line.

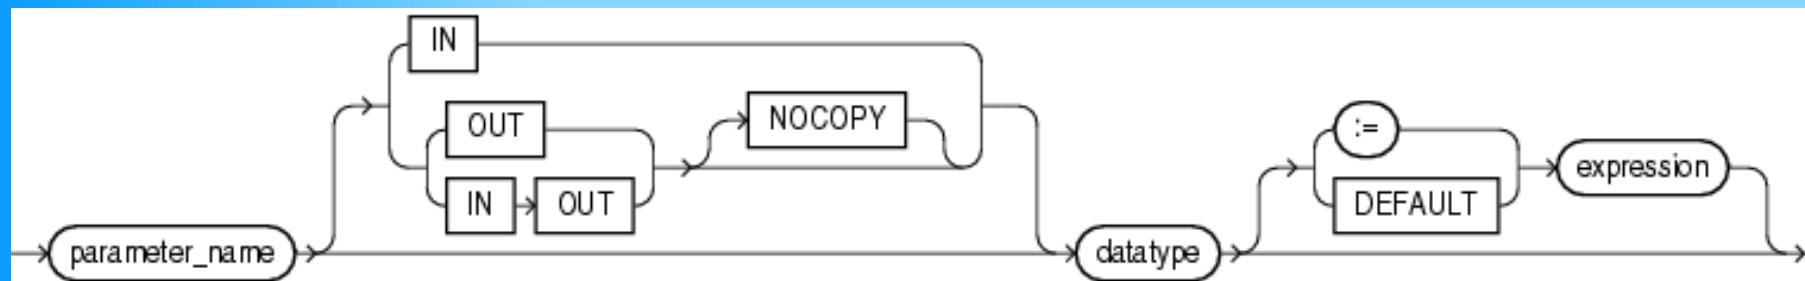
SQL:1999/IBM	MICROSOFT/SYBASE	ORACLE
CREATE PROCEDURE	CREATE PROCEDURE	CREATE PROCEDURE
Sp_proc1	Sp_proc1	Sp_proc1
(param1 INT)	@param1 INT	(param1 IN OUT INT)
MODIFIES SQL DATA BEGIN DECLARE num1 INT;	AS DECLARE @num1 INT	AS num1 INT; BEGIN
IF param1 <> 0	IF @param1 <> 0	IF param1 <> 0
THEN SET param1 = 1;	SELECT @param1 = 1;	THEN param1 :=1;
END IF		END IF;
UPDATE Table1 SET column1 = param1;	UPDATE Table1 SET column1 = @param1	UPDATE Table1 SET column1 = param1;
END		END

Source: Data from *SQL Performance Tuning* (Gulutzan and Pelzer, Addison-Wesley, 2002). Viewed at www.tdan.com/i023fe03.htm, June 6, 2007 (no longer available from this site).

Oracle Syntax: Create a Function



Oracle Syntax: parameter_declaration



Example: The CREATE FUNCTION Statement in Oracle

Create the function **get_bal** on table Orders_T to return the balance of a specified account:

```
CREATE FUNCTION get_bal(acc_no IN NUMBER)
RETURN NUMBER
IS acc_bal NUMBER(11,2);
BEGIN
    SELECT order_total
    INTO acc_bal
    FROM Orders_T
    WHERE customer_id = acc_no;
    RETURN(acc_bal);
END;
/
```

To get the balance of the account number 165:

```
SELECT get_bal(165) FROM DUAL;
```

GET_BAL(165)

2519

Example: The CREATE PROCEDURE Statement in Oracle

To update the salary of an employee on table Employees_T:

```
CREATE PROCEDURE Update_salary (em_id NUMBER, new_salary  
NUMBER) AS  
    updated boolean;  
BEGIN  
    UPDATE Employees_T  
    SET Salary=new_salary  
    WHERE Employees_T.EMPLOYEE_ID = em_id;  
    updated := true;  
END;  
/
```

Update the salary of employee # 100 to \$500,000:

EXECUTE Update_salary(100, 500000);

Example: The CREATE PROCEDURE Statement without Parameters in Oracle

To calculate a sale price based on the standard price of the product:

```
CREATE PROCEDURE ProductLineSale AS
BEGIN
    UPDATE Product_T
        SET SalePrice = 0.9*ProductStandardPrice
        WHERE ProductStandardPrice >= 400;
    UPDATE Product_T
        SET SalePrice = 0.85*ProductStandardPrice
        WHERE ProductStandardPrice < 400;
END;
/
```

To run the procedure:

EXECUTE ProductLineSale;

Note: if the SalePrice column does not exist, run the following statement:

**ALTER TABLE Product_T
ADD (SalePrice NUMBER (6,2));**

CREATE TRIGGER syntax

```
CREATE [OR REPLACE] TRIGGER <trigger_name>
[<ENABLE | DISABLE>]
<BEFORE | AFTER> <ACTION> [OR <ACTION> OR
<ACTION>]
ON <table_name>
[FOLLOWS <schema.trigger_name>]
[FOR EACH ROW]
DECLARE
    <variable definitions>
BEGIN
    <trigger_code>
EXCEPTION
    <exception clauses>
END <trigger_name>;
/
```

CREATE TRIGGER

```
CREATE OR REPLACE TRIGGER orders_after_update
AFTER UPDATE
  ON orders_t
  FOR EACH ROW
DECLARE
  v_username varchar2(10);
BEGIN
  -- Find username of a person and perform UPDATE into table
  SELECT user INTO v_username
  FROM dual;
  -- Insert record into an audit table created earlier
  INSERT INTO orders_audit_t
    ( order_id, quantity_before, quantity_after, username )
  VALUES
    ( :new.order_id, :old.quantity, :new.quantity, v_username );
END;
/
```

Another Example: CREATE TRIGGER

```
CREATE OR REPLACE TRIGGER xy.salary_check
BEFORE INSERT OR UPDATE OF salary, jobtitle ON xy.emp_t
FOR EACH ROW
BEGIN
    WHEN (new.jobtitle <> 'Data Scientist')
        CALL check_salary(:new.jobtitle, :new.salary, :new.empname);
END;
/
```

Note: The `check_salary()` procedure has been implemented earlier.

Embedded and Dynamic SQL

- Embedded SQL
 - Including hard-coded SQL statements in a program written in another language such as C or Java
- Dynamic SQL
 - Ability for an application program to generate SQL code on the fly, as the application is running

Reasons to Embed SQL in 3GL (3rd Generation Languages)

- Can create a more flexible, accessible interface for the users.
- Possible performance improvement.
- Database security improvement; grant access only to the application instead of users.

Example: Embedded SQL in C

```
#include <stdio.h>

/* declare host variables */
char userid[12]="SCOTT/TIGER";
char emp_name[10];
void sql_error();

/* include the SQL Communications Area */
#include<sqlca.h>

void main()
{
/* handle errors */
    EXEC SQL WHENEVER SQLERROR do sql_error("Oracle
error");

/* connect to Oracle */
    EXEC SQL CONNECT :userid;
    printf("Connected.\n");

/* declare a scrollable cursor */
    EXEC SQL DECLARE emp_cursor SCROLL CURSOR FOR
        SELECT ename FROM emp;

/* open the cursor and identify the active set */
    EXEC SQL OPEN emp_cursor;

/* Fetch the last row */
    EXEC SQL FETCH LAST emp_cursor INTO
        :emp_name;

/* Fetch row number 5 */
    EXEC SQL FETCH ABSOLUTE 5 emp_cursor INTO
        :emp_name;

/* Fetch row number 10 */
    EXEC SQL FETCH RELATIVE 5 emp_cursor INTO
        :emp_name;

/* Fetch row number 7 */
    EXEC SQL FETCH RELATIVE -3 emp_cursor INTO
        :emp_name;

/* Fetch the first row */
    EXEC SQL FETCH FIRST emp_cursor INTO
        :emp_name;

/* Fetch row number 2*/
    EXEC SQL FETCH my_cursor INTO :emp_name;

/* Fetch row number 3 */
    EXEC SQL FETCH NEXT my_cursor INTO :emp_name;

/* Fetch row number 3 */
    EXEC SQL FETCH CURRENT my_cursor INTO
        :emp_name;

/* Fetch row number 2 */
    EXEC SQL FETCH PRIOR my_cursor INTO
        :emp_name;
}
```