

Machine Learning for Data Science (CS4786)

Lecture 8

Kernel PCA
&
Isomap + TSNE

LINEAR PROJECTIONS

$$n \begin{matrix} X \\ \times d \\ W \end{matrix} = n \begin{matrix} Y \\ K \end{matrix}$$

d K

Works when data lies in a low dimensional linear sub-space

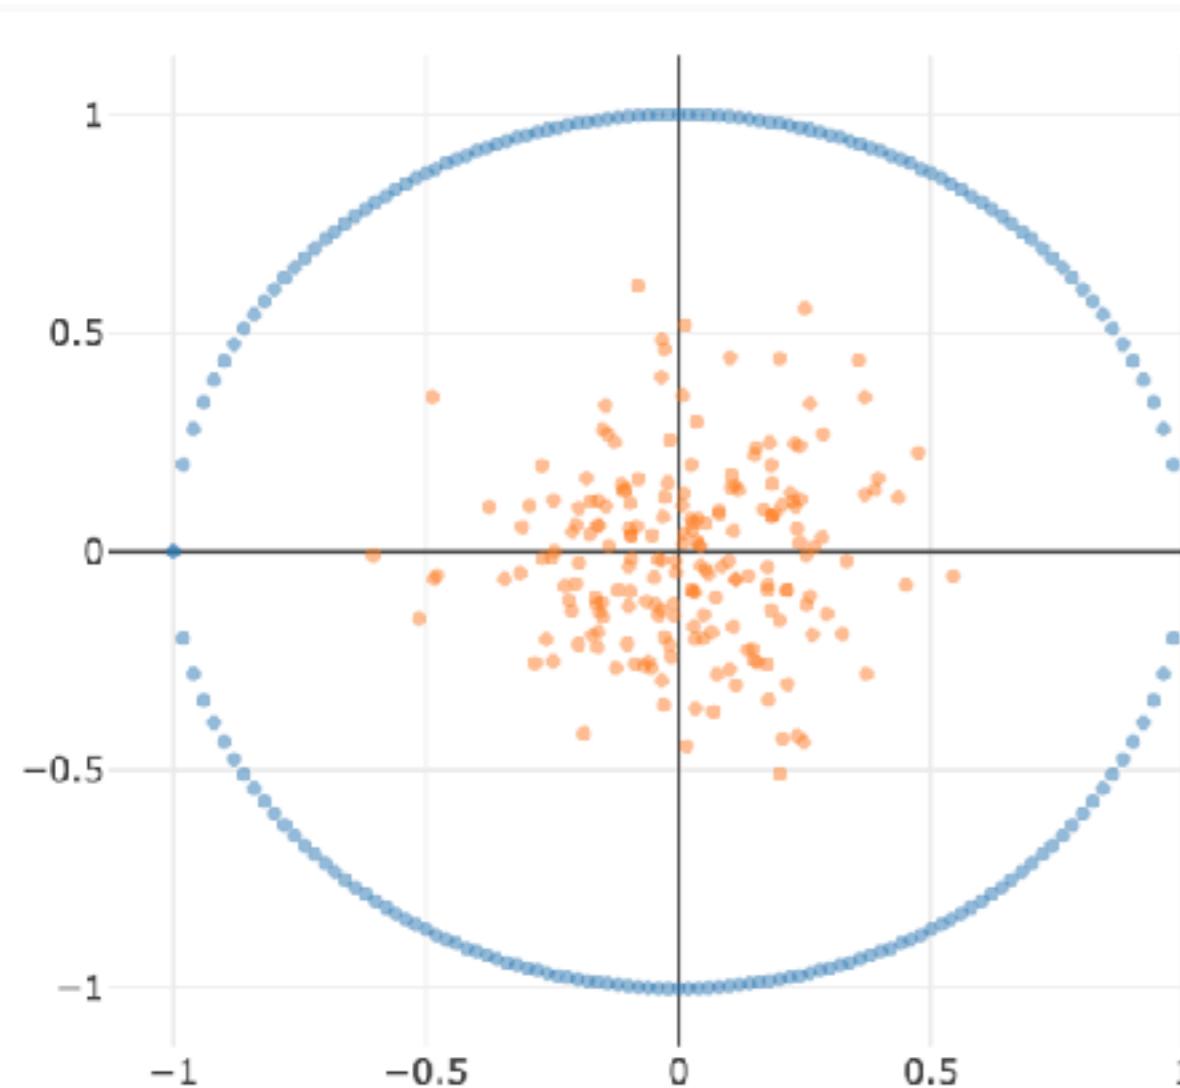
KERNEL TRICK

- We have nice methods for linear dimensionality reduction
- Can we use this beyond the linear realm?

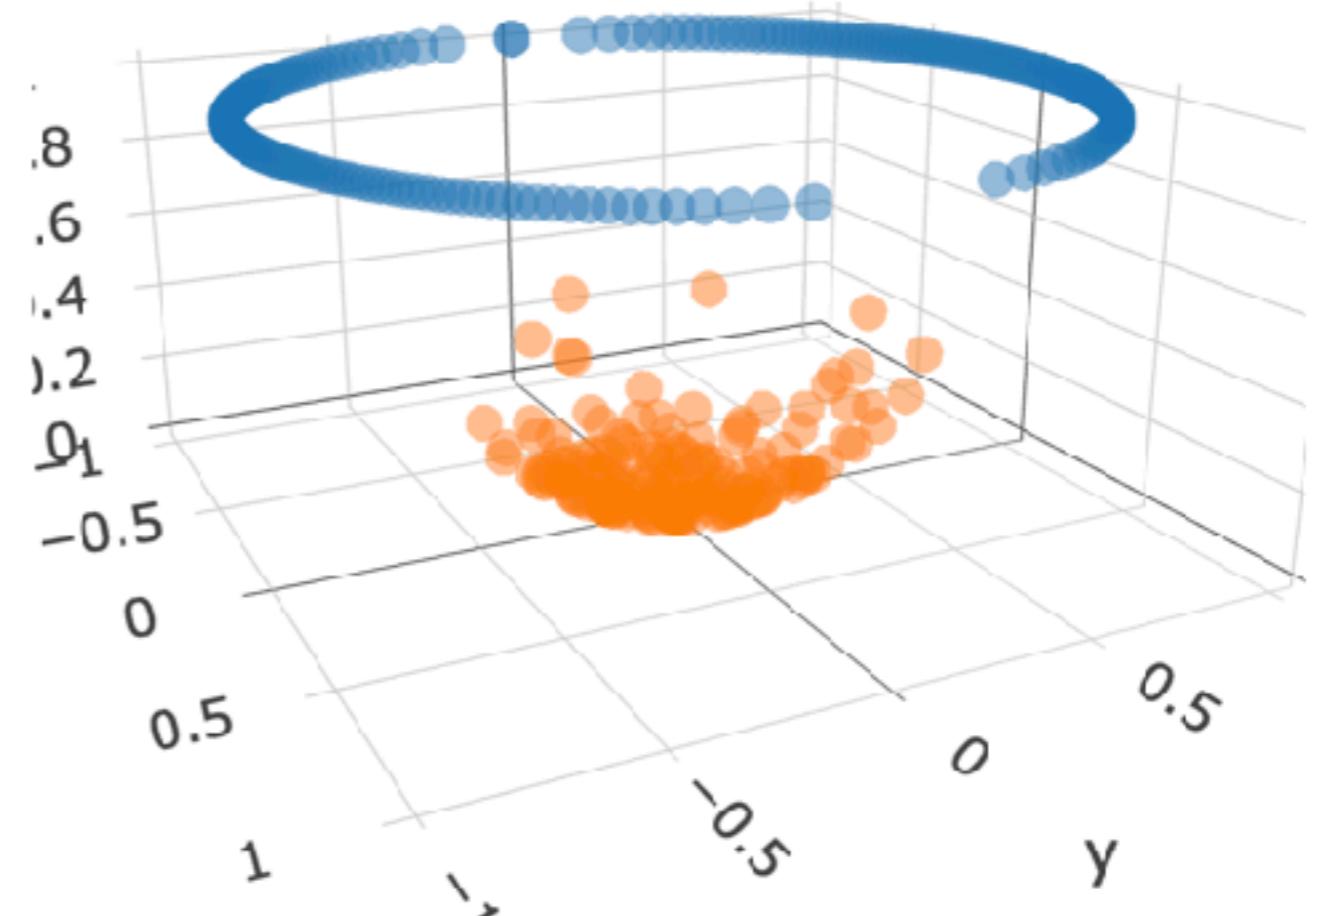
KERNEL TRICK

- Lift to higher dimensions (introduces non-linearity)
- Perform linear dimensionality reduction in this high dimensional space

EXAMPLE



Original Data in 2D
 (x, y)



Data Lifted to 3D
 $(x, y, x^2 + y^2)$

A FIRST CUT

- Given $\mathbf{x}_t \in \mathbb{R}^d$, the feature space vector is given by mapping

$$\Phi(\mathbf{x}_t) = (\mathbf{x}_t[1], \dots, \mathbf{x}_t[d], \mathbf{x}_t[1] \cdot \mathbf{x}_t[1], \mathbf{x}_t[1] \cdot \mathbf{x}_t[2], \dots, \mathbf{x}_t[d] \cdot \mathbf{x}_t[d], \dots)^{\top}$$

- Enumerating products up to order K (ie. products of at most K coordinates) we can get degree K polynomials.
- However dimension blows up as d^K
- Is there a way to do this without enumerating Φ ?

KERNEL TRICK

- Essence of Kernel trick:
 - If we can write down an algorithm only in terms of $\Phi(\mathbf{x}_t)^\top \Phi(\mathbf{x}_s)$ for data points \mathbf{x}_t and \mathbf{x}_s
 - Then we don't need to explicitly enumerate $\Phi(\mathbf{x}_t)$'s but instead, compute $k(\mathbf{x}_t, \mathbf{x}_s) = \Phi(\mathbf{x}_t)^\top \Phi(\mathbf{x}_s)$ (even if Φ maps to infinite dimensional space)
- Example: RBF kernel $k(\mathbf{x}_t, \mathbf{x}_s) = \exp(-\sigma \|\mathbf{x}_t - \mathbf{x}_s\|_2^2)$, polynomial kernel $k(\mathbf{x}_t, \mathbf{x}_s) = (\mathbf{x}_t^\top \mathbf{y}_t)^p$
- Kernel function measures similarity between points.

KERNEL TRICK

$$\begin{aligned} (\mathbf{x}_t^\top \mathbf{y}_t)^p &= \sum_{k_1+k_2+\dots+k_d=p} \binom{p}{k_1, k_2, \dots, k_d} \prod_{j=1}^d (x_t[j] y_t[j])^{k_j} \\ &= \sum_{k_1+k_2+\dots+k_d=p} \left(\sqrt{\binom{p}{k_1, k_2, \dots, k_d}} \prod_{j=1}^d x_t[j]^{k_j} \right) \cdot \left(\sqrt{\binom{p}{k_1, k_2, \dots, k_d}} \prod_{j=1}^d y_t[j]^{k_j} \right) \\ \Phi(\mathbf{x})^\top &= \left(\dots, \sqrt{\binom{p}{k_1, k_2, \dots, k_d}} \prod_{j=1}^d x_t[j]^{k_j}, \dots \right)_{k_1+k_2+\dots+k_d=p} \end{aligned}$$

Key Idea:

If an algorithm only depends on inner products,
we can simply replace inner product in x space
by inner product in $\phi(x)$ space

Can we write PCA so it only depends on inner products?

LETS REWRITE PCA

Lets start with the assumption that Data is centered! (i.e. Sum of \mathbf{x}_t 's is 0)

- k^{th} column of \mathbf{W} is eigenvector of covariance matrix
That is, $\lambda_k \mathbf{W}_k = \Sigma \mathbf{W}_k$. Rewriting, for centered \mathbf{X}

$$\lambda_k \mathbf{W}_k = \frac{1}{n} \left(\sum_{t=1}^n \mathbf{x}_t \mathbf{x}_t^\top \right) \mathbf{W}_k = \frac{1}{n} \sum_{t=1}^n (\mathbf{x}_t^\top \mathbf{W}_k) \mathbf{x}_t$$

But $\mathbf{x}_t^\top \mathbf{W}_k = \mathbf{y}_t[k]$

$$\lambda_k \mathbf{W}_k = \frac{1}{n} \sum_{t=1}^n \mathbf{y}_t[k] \mathbf{x}_t$$

LETS REWRITE PCA

$$\begin{aligned}\mathbf{y}_s[k] &= W_k^\top \mathbf{x}_s \\ &= \frac{1}{\lambda_k} \left(\frac{1}{n} \sum_{t=1}^n \mathbf{y}_t[k] \mathbf{x}_t \right)^\top \mathbf{x}_s \\ &= \frac{1}{n\lambda_k} \sum_{t=1}^n \mathbf{y}_t[k] \mathbf{x}_t^\top \mathbf{x}_s \\ &= \frac{1}{n\lambda_k} \sum_{t=1}^n \mathbf{y}_t[k] \tilde{K}_{s,t}\end{aligned}$$

Where $\tilde{K}_{s,t} = \mathbf{x}_t^\top \mathbf{x}_s$ is the kernel matrix for centered data

LETS REWRITE PCA

- Hence, the k'th column on Y matrix is such that

$$\mathbf{y}[k] = \frac{1}{n\lambda_k} \mathbf{y}[k] \tilde{K}$$

Also we have, $1 = \|W_k\|^2 = \frac{1}{\lambda_k^2 n^2} \left(\sum_{t=1}^n \mathbf{y}_t[k] \mathbf{x}_t \right)^\top \left(\sum_{s=1}^n \mathbf{y}_s[k] \mathbf{x}_s \right)$

$$= \frac{1}{\lambda_k^2 n^2} \sum_{t=1}^n \sum_{s=1}^n \mathbf{y}_s[k] \mathbf{x}_s^\top \mathbf{x}_t \mathbf{y}_t[k]$$

$$= \frac{1}{\lambda_k^2 n^2} \mathbf{y}[k] \tilde{K} \mathbf{y}[k]^\top = \frac{1}{n\lambda_k} \|\mathbf{y}[k]\|^2$$

Hence $P_k = \mathbf{y}[k]/\sqrt{n\lambda_k}$ is an eigenvector of \tilde{K} with eigen value $\gamma_k = n\lambda_k$

REWRITTING PCA

- We assumed centered data, what if its not,

$$\begin{aligned}\tilde{K}_{s,t} &= \left(\mathbf{x}_t - \frac{1}{n} \sum_{u=1}^n \mathbf{x}_u \right)^\top \left(\mathbf{x}_s - \frac{1}{n} \sum_{u=1}^n \mathbf{x}_u \right) \\ &= \mathbf{x}_t^\top \mathbf{x}_s - \left(\frac{1}{n} \sum_{u=1}^n \mathbf{x}_u \right)^\top \mathbf{x}_s - \left(\frac{1}{n} \sum_{u=1}^n \mathbf{x}_u \right)^\top \mathbf{x}_t \\ &\quad + \frac{1}{n^2} \left(\sum_{u=1}^n \mathbf{x}_u \right)^\top \left(\sum_{v=1}^n \mathbf{x}_v \right) \\ &= \mathbf{x}_t^\top \mathbf{x}_s - \frac{1}{n} \sum_{u=1}^n \mathbf{x}_u^\top \mathbf{x}_s - \frac{1}{n} \sum_{u=1}^n \mathbf{x}_u^\top \mathbf{x}_t + \frac{1}{n^2} \sum_{u=1}^n \sum_{v=1}^n \mathbf{x}_u^\top \mathbf{x}_v\end{aligned}$$

REWRITING PCA

- Equivalently, if Kern is the matrix ($\text{Kern}_{t,s} = \mathbf{x}_t^\top \mathbf{x}_s$),

$$\tilde{\mathbf{K}} = \mathbf{Kern} - \frac{(\mathbf{1}_{n \times n} \times \mathbf{Kern})}{n} - \frac{(\mathbf{Kern} \times \mathbf{1}_{n \times n})}{n} + \frac{(\mathbf{1}_{n \times n} \times \mathbf{Kern} \times \mathbf{1}_{n \times n})}{n^2}$$

KERNEL PCA

1.

n
Kern

n

$$\text{Kern} = \begin{bmatrix} k(x_1, x_1) & k(x_1, x_2) & \dots & k(x_1, x_n) \\ k(x_2, x_1) & k(x_2, x_2) & \dots & k(x_2, x_n) \\ \vdots & \vdots & \ddots & \vdots \\ \vdots & \vdots & \vdots & \vdots \\ k(x_{n-1}, x_1) & k(x_{n-1}, x_2) & \dots & k(x_{n-1}, x_n) \\ k(x_n, x_1) & k(x_n, x_2) & \dots & k(x_n, x_n) \end{bmatrix}$$

2.

n
 \tilde{K}

$$\tilde{K} = \text{Kern} - \frac{1}{n} (\mathbf{1} \text{ Kern} + \text{Kern} \mathbf{1}) + \frac{1}{n^2} \mathbf{1} \text{ Kern} \mathbf{1}$$

n

KERNEL PCA

$$3. \left[\begin{smallmatrix} n & P \\ K & \gamma \end{smallmatrix} \right] = \text{eigs}\left(\begin{smallmatrix} \tilde{K} \\ , K \end{smallmatrix} \right)$$

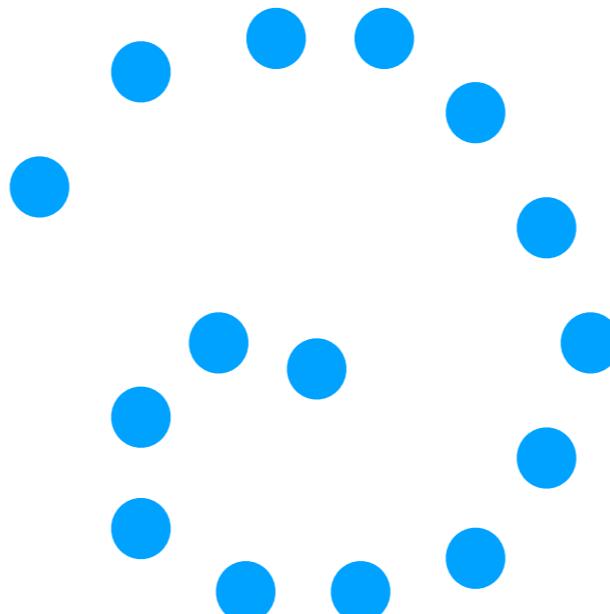
$$4. \begin{smallmatrix} n & Y \\ K & \end{smallmatrix} = \begin{smallmatrix} & & \\ \vdots & & \vdots \\ P_1\sqrt{\gamma_1} & & P_K\sqrt{\gamma_K} \\ \vdots & & \vdots \\ & & \end{smallmatrix}$$

Kernel Methods: A note

- We can kernelize CCA and any other linear dimensionality reduction method.
- For any linear method, solution lies within linear span of data
- Hence y 's can be computed only based on inner products.

MANIFOLD BASED DIMENSIONALITY REDUCTION

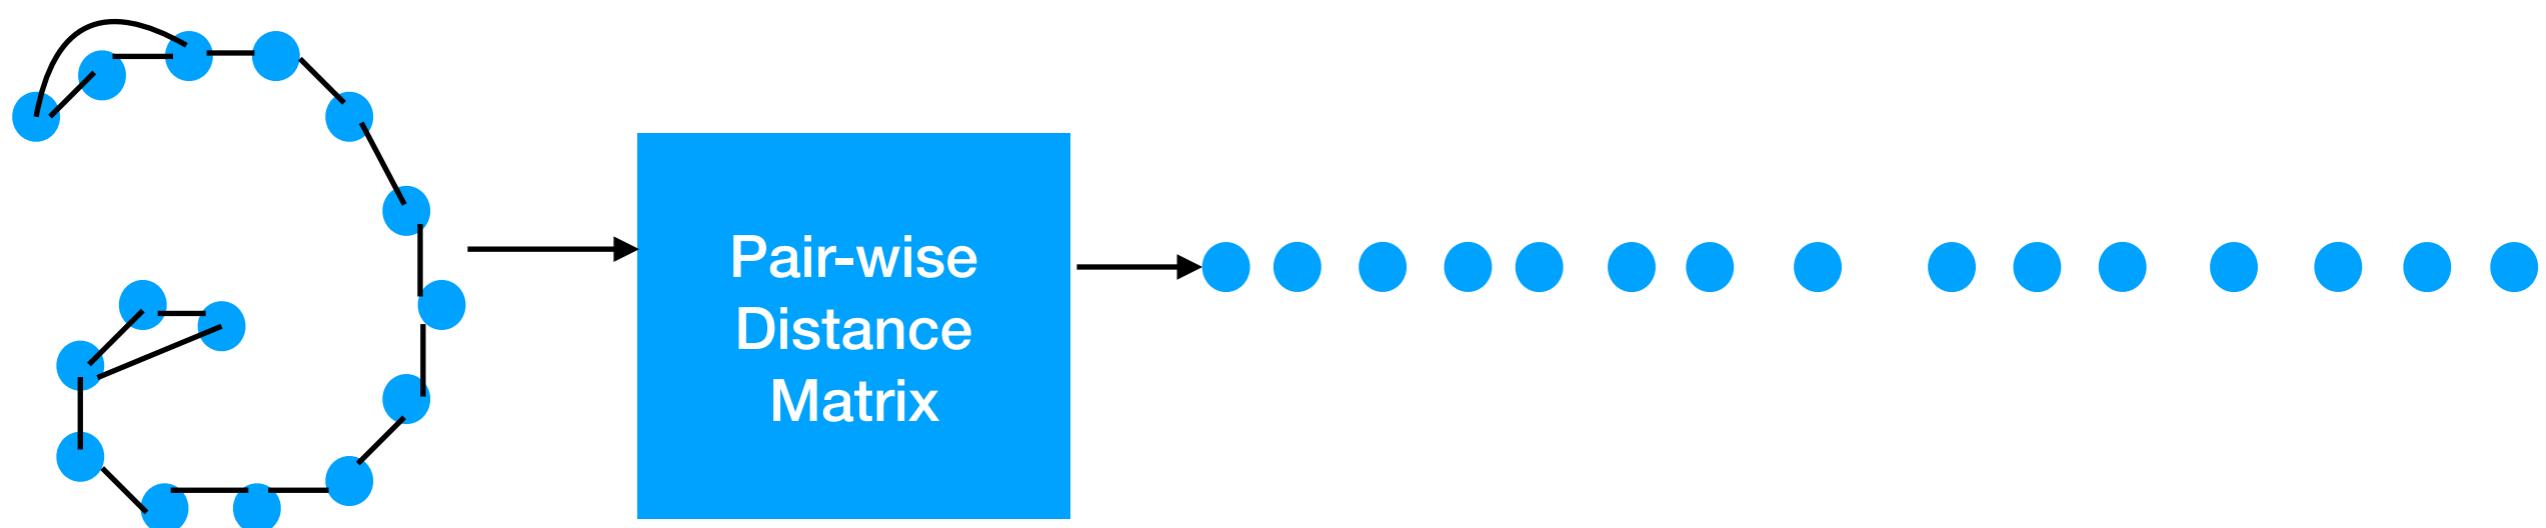
- Key Assumption: Points live on a low dimensional manifold
- Manifold: subspace that looks locally Euclidean
- Given data, can we uncover this manifold?



Can we unfold this?

METHOD I: ISOMAP

- ① For every point, find its (k -) Nearest Neighbors
- ② Form the Nearest Neighbor graph
- ③ For every pair of points A and B , distance between point A to B is shortest distance between A and B on graph
- ④ Find points in low dimensional space such that distances between points in this space is equal to distance on graph.



ISOMAP: PITFALLS

- ① If we don't take enough nearest neighbors, then graph may not be connected
- ② If we connect points too far away, points that should not be connected can get connected
- ③ There may not be a right number of nearest neighbors we should consider!

STOCHASTIC NEIGHBORHOOD EMBEDDING

- Use a probabilistic notion of which points are neighbors.
- Close by points are neighbors with high probability, . . .
Eg: For point \mathbf{x}_t , point \mathbf{x}_s is picked as neighbor with probability

$$p_{t \rightarrow s} = \frac{\exp\left(-\frac{\|\mathbf{x}_s - \mathbf{x}_t\|^2}{2\sigma^2}\right)}{\sum_{u \neq t} \exp\left(-\frac{\|\mathbf{x}_u - \mathbf{x}_t\|^2}{2\sigma^2}\right)}$$

Probability that points s and t are connected $P_{s,t} = P_{t,s} = \frac{p_{t \rightarrow s} + p_{s \rightarrow t}}{2n}$

- Goal: Find $\mathbf{y}_1, \dots, \mathbf{y}_n$ with stochastic neighborhood distribution Q such that “ P and Q are similar”

i.e. minimize:

$$\text{KL}(P \| Q) = \sum_{s,t} P_{s,t} \log \left(\frac{P_{s,t}}{Q_{s,t}} \right) = \sum_{s,t} P_{s,t} \log (P_{s,t}) - \sum_{s,t} P_{s,t} \log (Q_{s,t})$$

CHOICE FOR Q

- Just like we defined P , we can define Q for a given $\mathbf{y}_1, \dots, \mathbf{y}_n$ by

$$q_{t \rightarrow s} = \frac{\exp\left(-\frac{\|\mathbf{y}_s - \mathbf{y}_t\|^2}{2\sigma^2}\right)}{\sum_{u \neq t} \exp\left(-\frac{\|\mathbf{y}_u - \mathbf{y}_t\|^2}{2\sigma^2}\right)}$$

and then set $Q_{s,t} = \frac{q_{t \rightarrow s} + q_{s \rightarrow t}}{2n}$

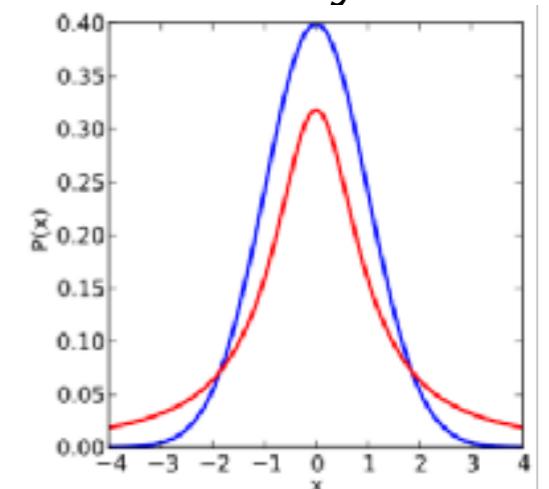
- However we are faced with the crowding problem:
 - In high dimension we have a lot of space, Eg. in d dimension we have $d+1$ equidistant point
 - For d dimensional gaussians, most points are found at distance \sqrt{d} from mean!
 - If we use gaussians in both high and low dimensional space, all the points are squished in to a small space
 - Too many points crowd the center!

METHOD II: T-SNE

- Instead for Q we use, student t distribution which is heavy tailed:

$$q_{t \rightarrow s} = \frac{(1 + \|\mathbf{y}_s - \mathbf{y}_t\|^2)^{-1}}{\sum_{u \neq t} (1 + \|\mathbf{y}_u - \mathbf{y}_t\|^2)^{-1}}$$

and then set $Q_{s,t} = \frac{q_{t \rightarrow s} + q_{s \rightarrow t}}{2n}$



- It can be verified that

$$\nabla_{\mathbf{y}_t} \text{KL}(P \| Q) = 4 \sum_{s=1}^n (P_{s,t} - Q_{s,t}) (\mathbf{y}_t - \mathbf{y}_s) (1 + \|\mathbf{y}_s - \mathbf{y}_t\|^2)^{-1}$$

- Algorithm: Find $\mathbf{y}_1, \dots, \mathbf{y}_n$ by performing gradient descent